



The Ottawa
Hospital

RESEARCH
INSTITUTE

L'Hôpital
d'Ottawa

INSTITUT DE
RECHERCHE



uOttawa

Institut de recherche
sur le cerveau

Brain and Mind
Research Institute

Workshop on Applied Deep Learning in Intracranial Neurophysiology

Part 5 – Introduction to Recurrent Neural Networks
June 21, 2019

Presented by Chadwick Boulay, MSc, PhD
Sachs Lab

Superior arm-movement decoding from cortex with a new, unsupervised-learning algorithm

Dataset

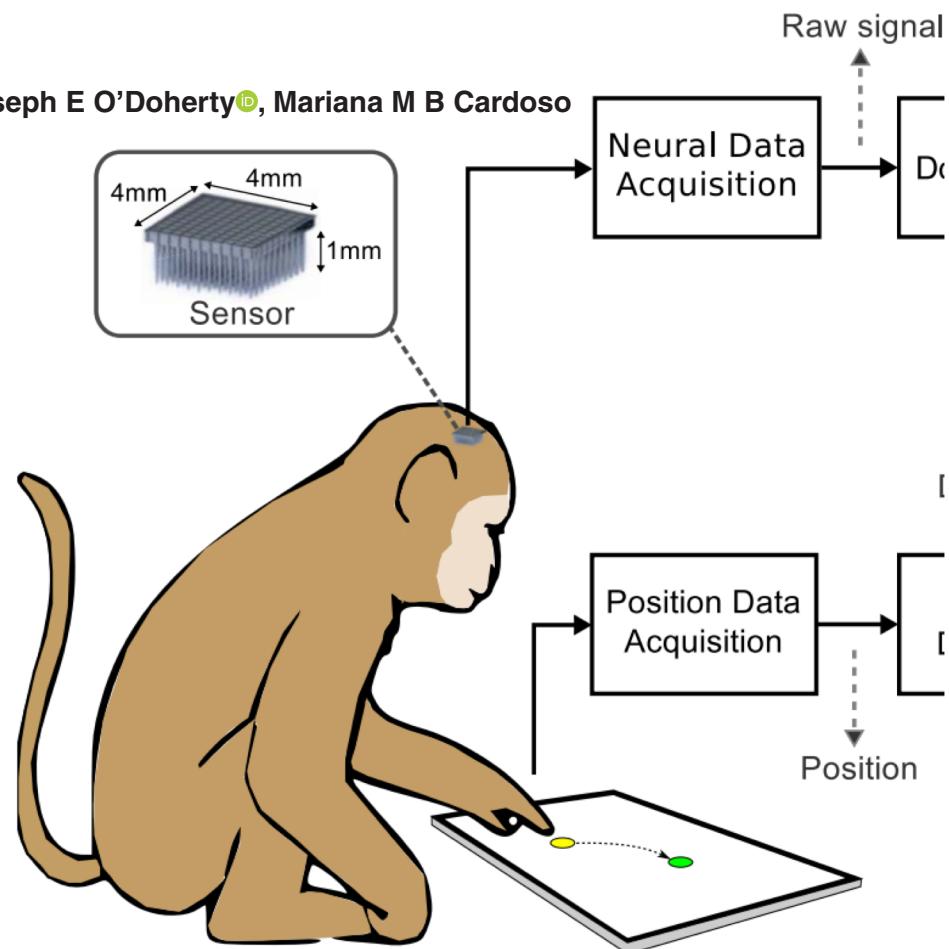
Joseph G Makin¹ , Joseph E O'Doherty , Mariana M B Cardoso and Philip N Sabes 

2 96-ch Utah arrays: M1 and S1

- 24.4 kHz
- BP filt 500-5000 Hz
- Thresholded at 3.5-4.0 STDs

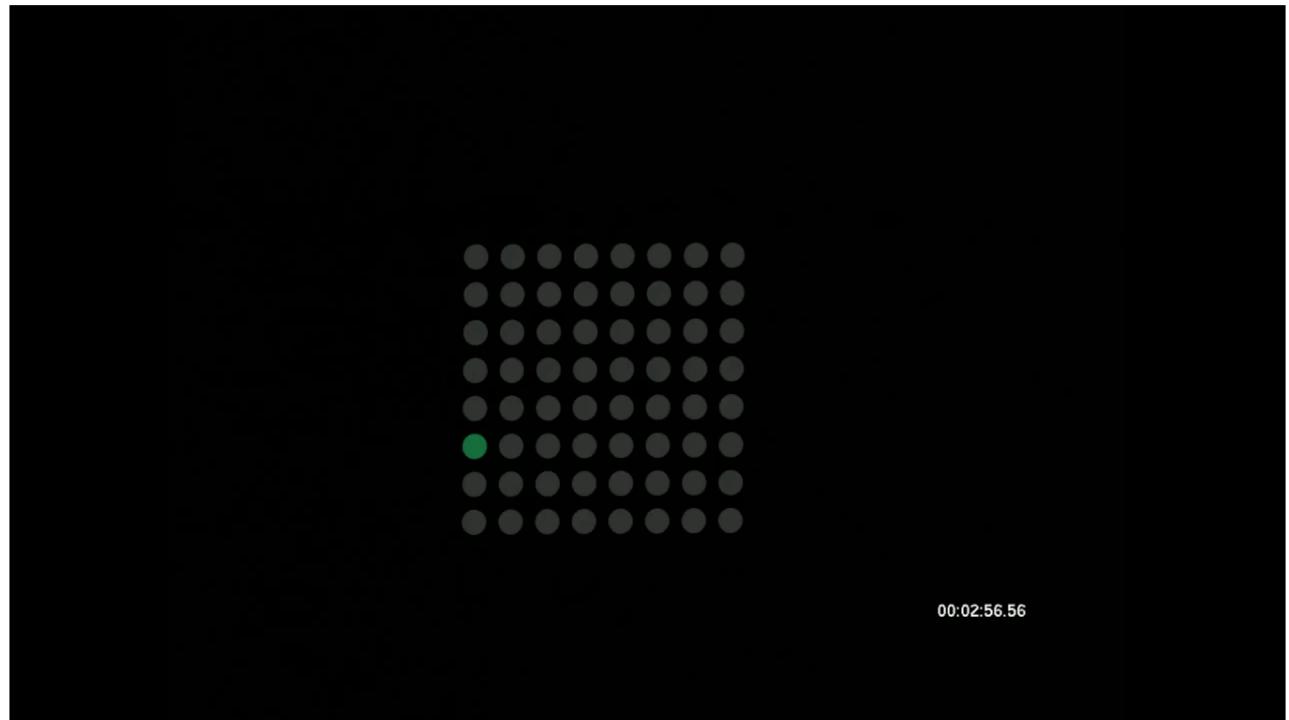
XYZ finger position recorded with Polhemus (magnetic) at 250 Hz

[Available online](#)



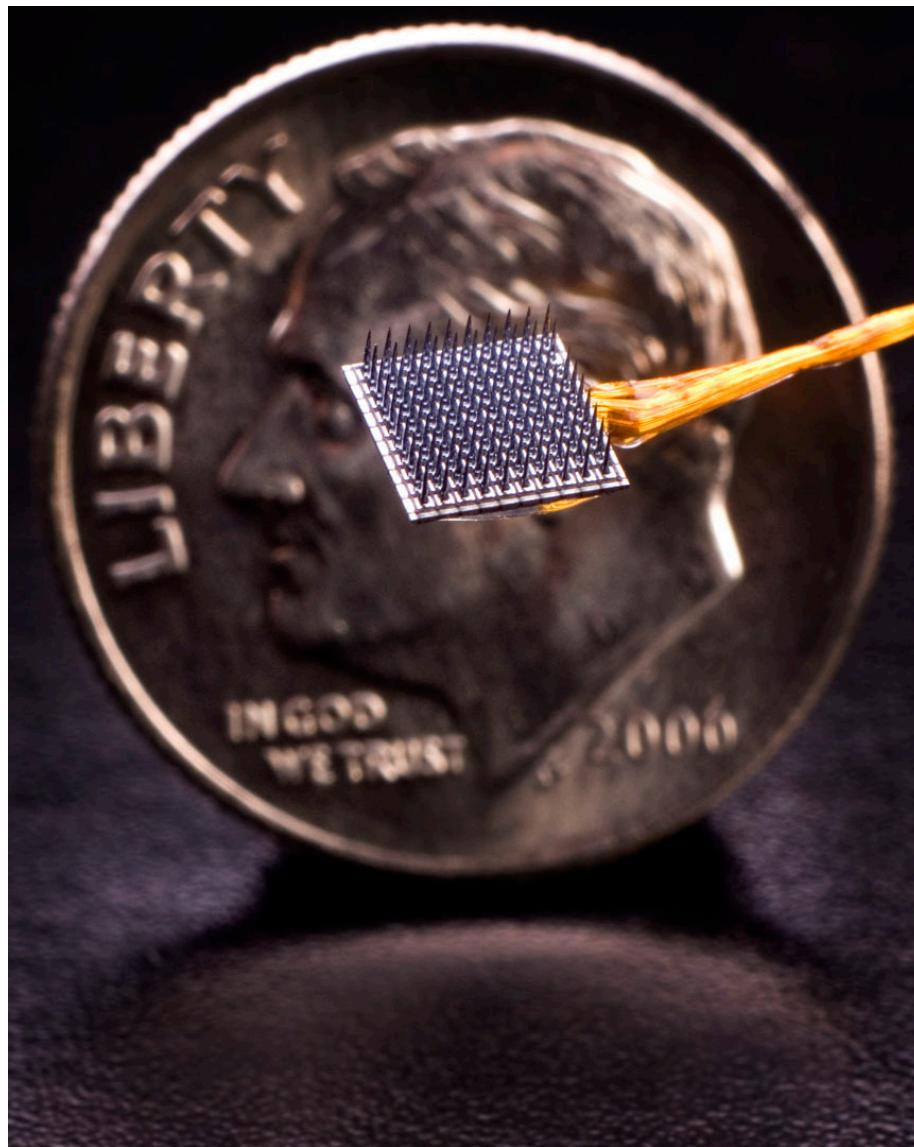
3D finger position
projected onto 2D
plane

$$\begin{pmatrix} 0 & 0 \\ -10 & 0 \\ 0 & -10 \end{pmatrix}$$



00:02:56.56

Utah Array





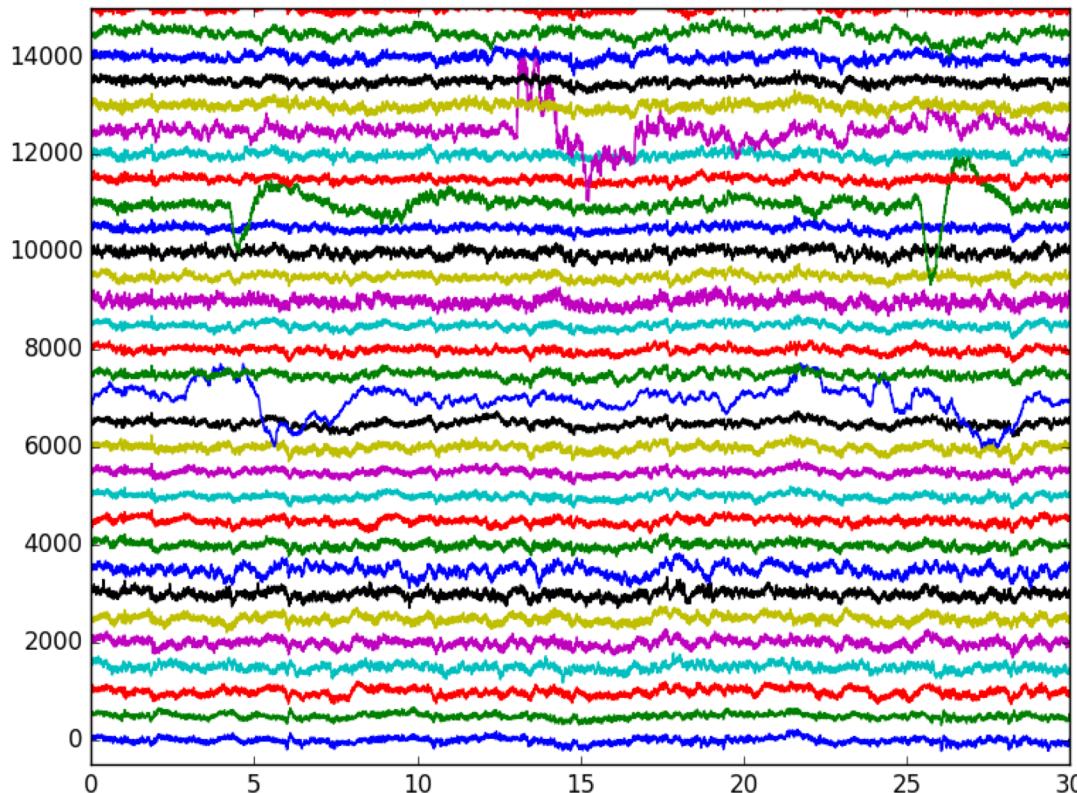
- Data contents:

- Events

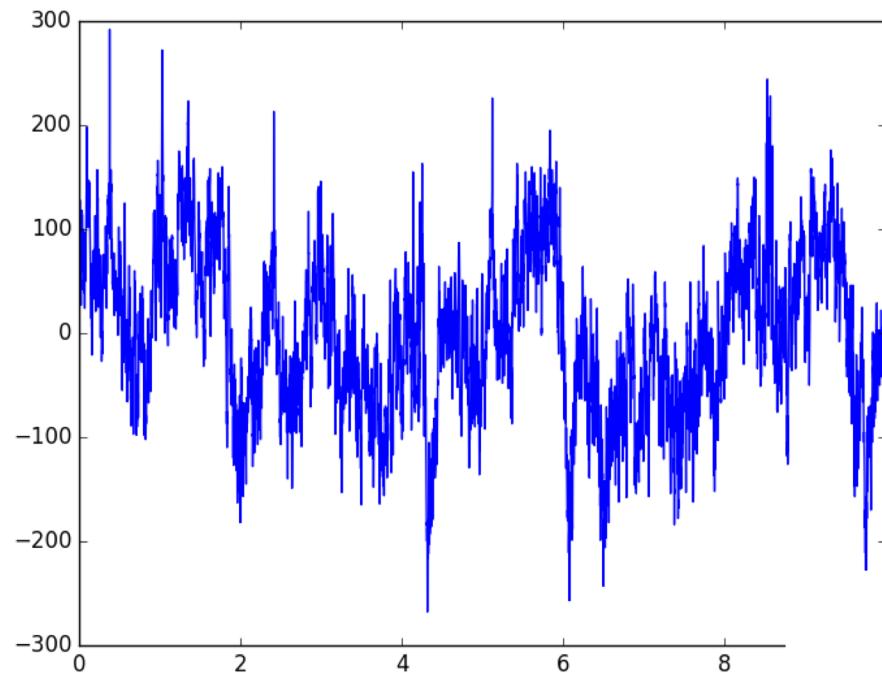
```
events = [(timestamp, event_string), ...]
```

- Broadband Field Potentials

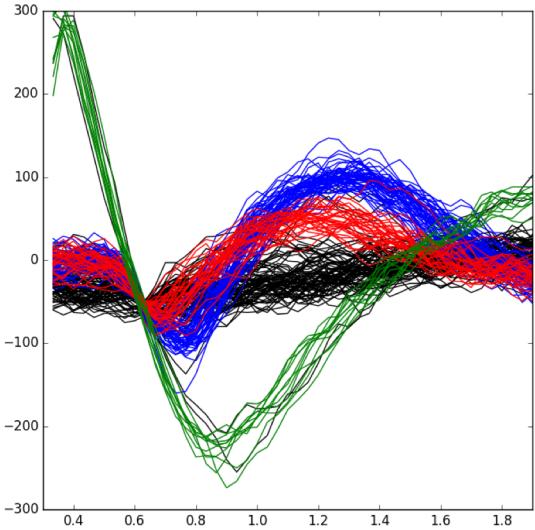
- Low-pass filter
 - Downsample (decimate) to get LFP
- High-pass filter
 - Get threshold crossing
 - Get spike waveforms



Raw

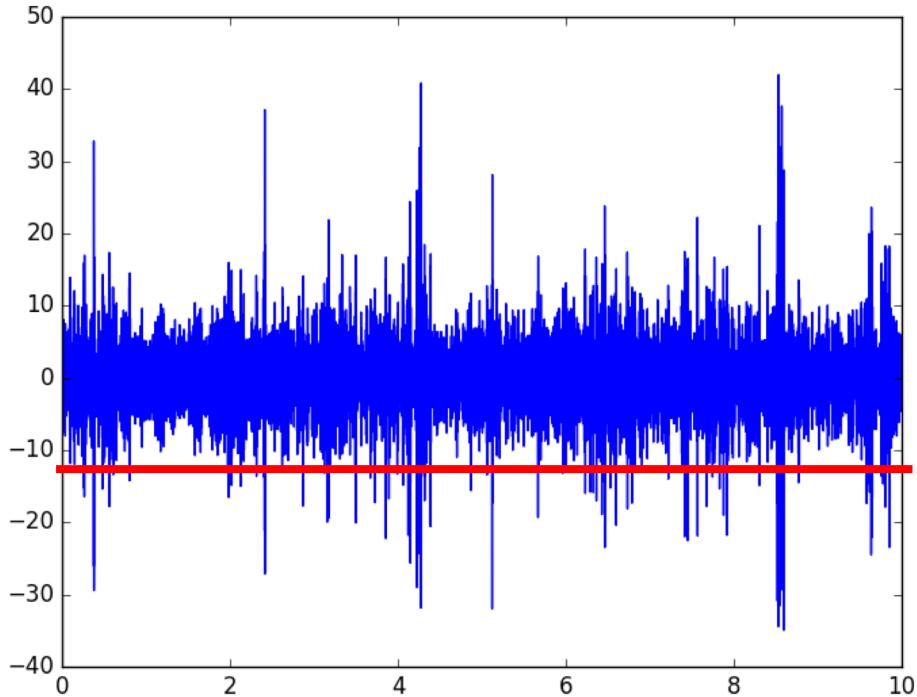


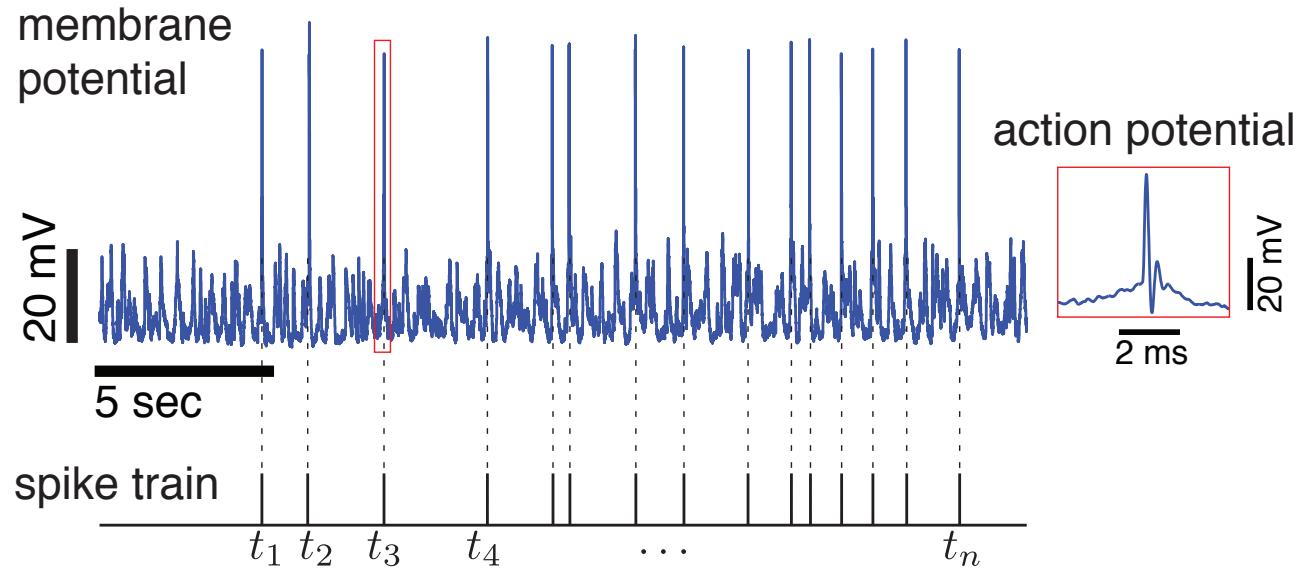
Preprocess



$\sim -3.5 - 4$ SD

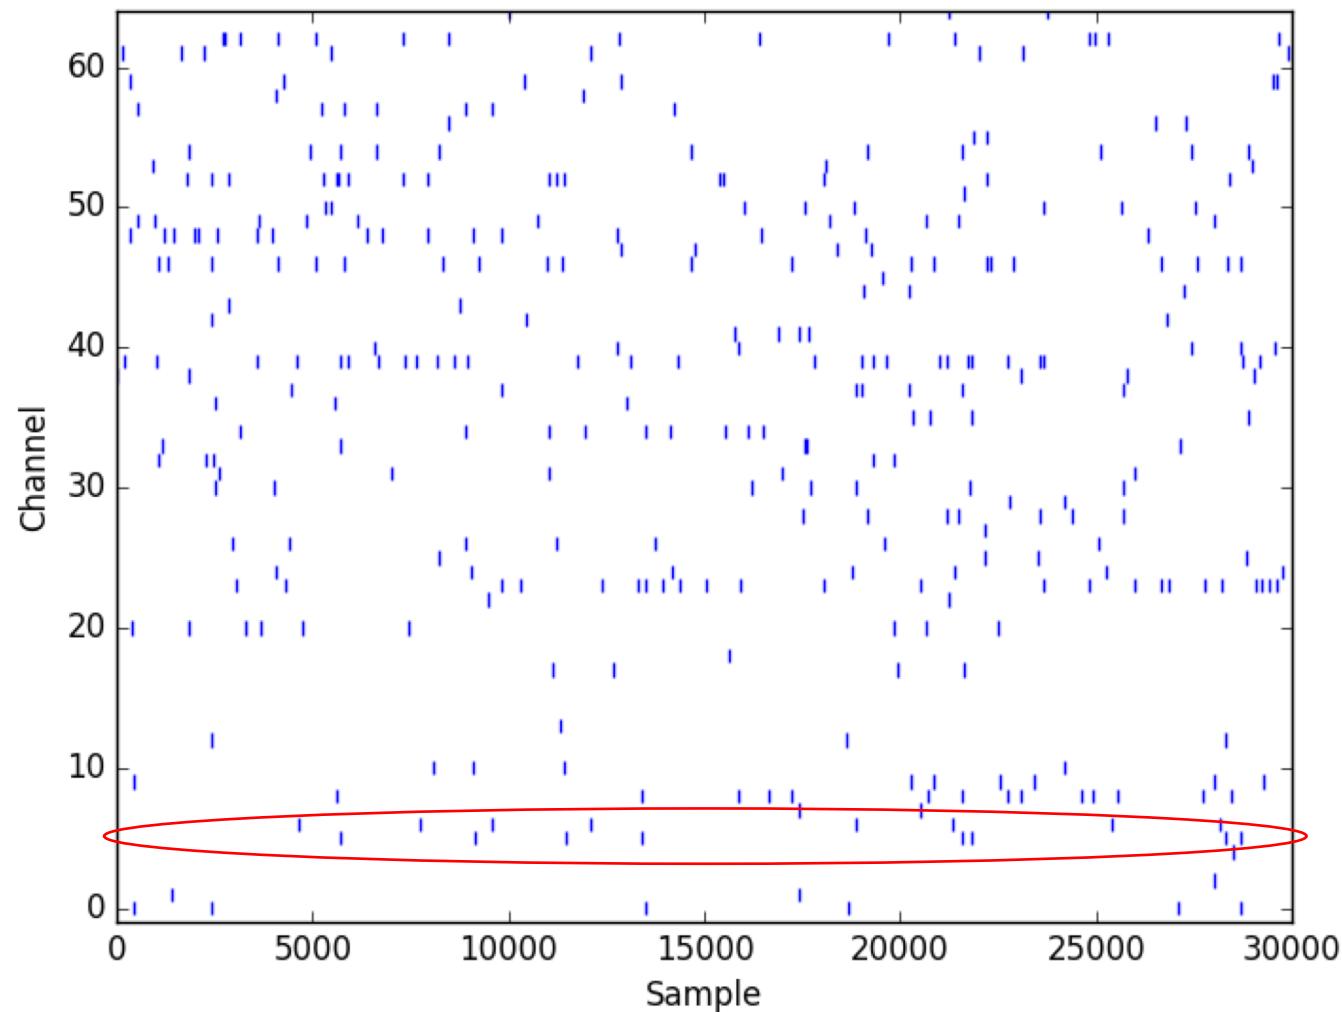
Highpass Filtered





Park et al., IEEE Sig. Proc. Mag. 2013

Spike Trains



Rate coding – Binned spike counts

Non-binned representation



High precision binning

0 1 0 0 0 1 0 0 1 0 1 0 0 0 0 0

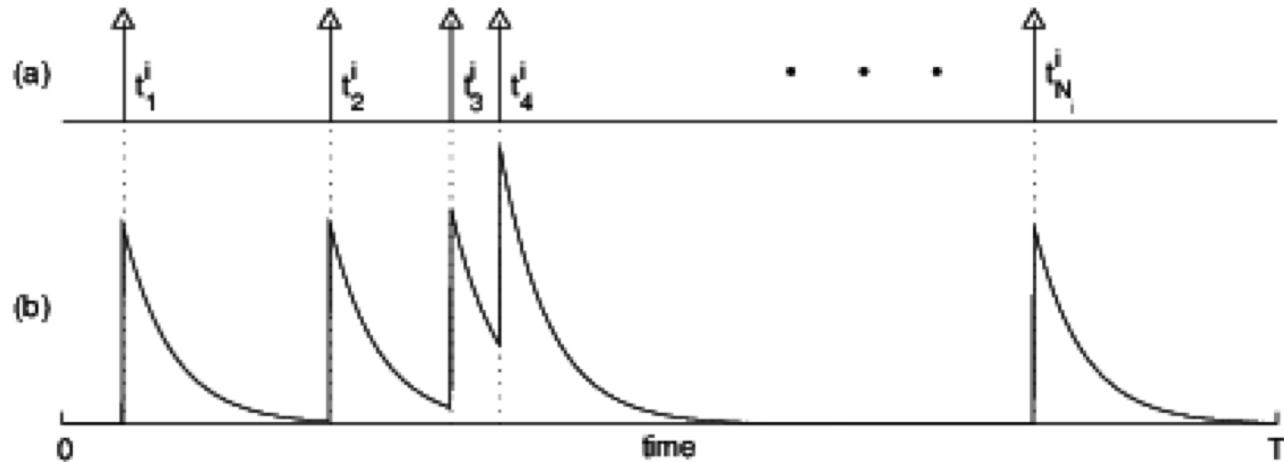
0 1 0 0 1 0 0 0 1 1 0 0 0 0 0 0

Low precision binning

1 1 2 0

1 1 2 0

Rate coding – Kernel convolution



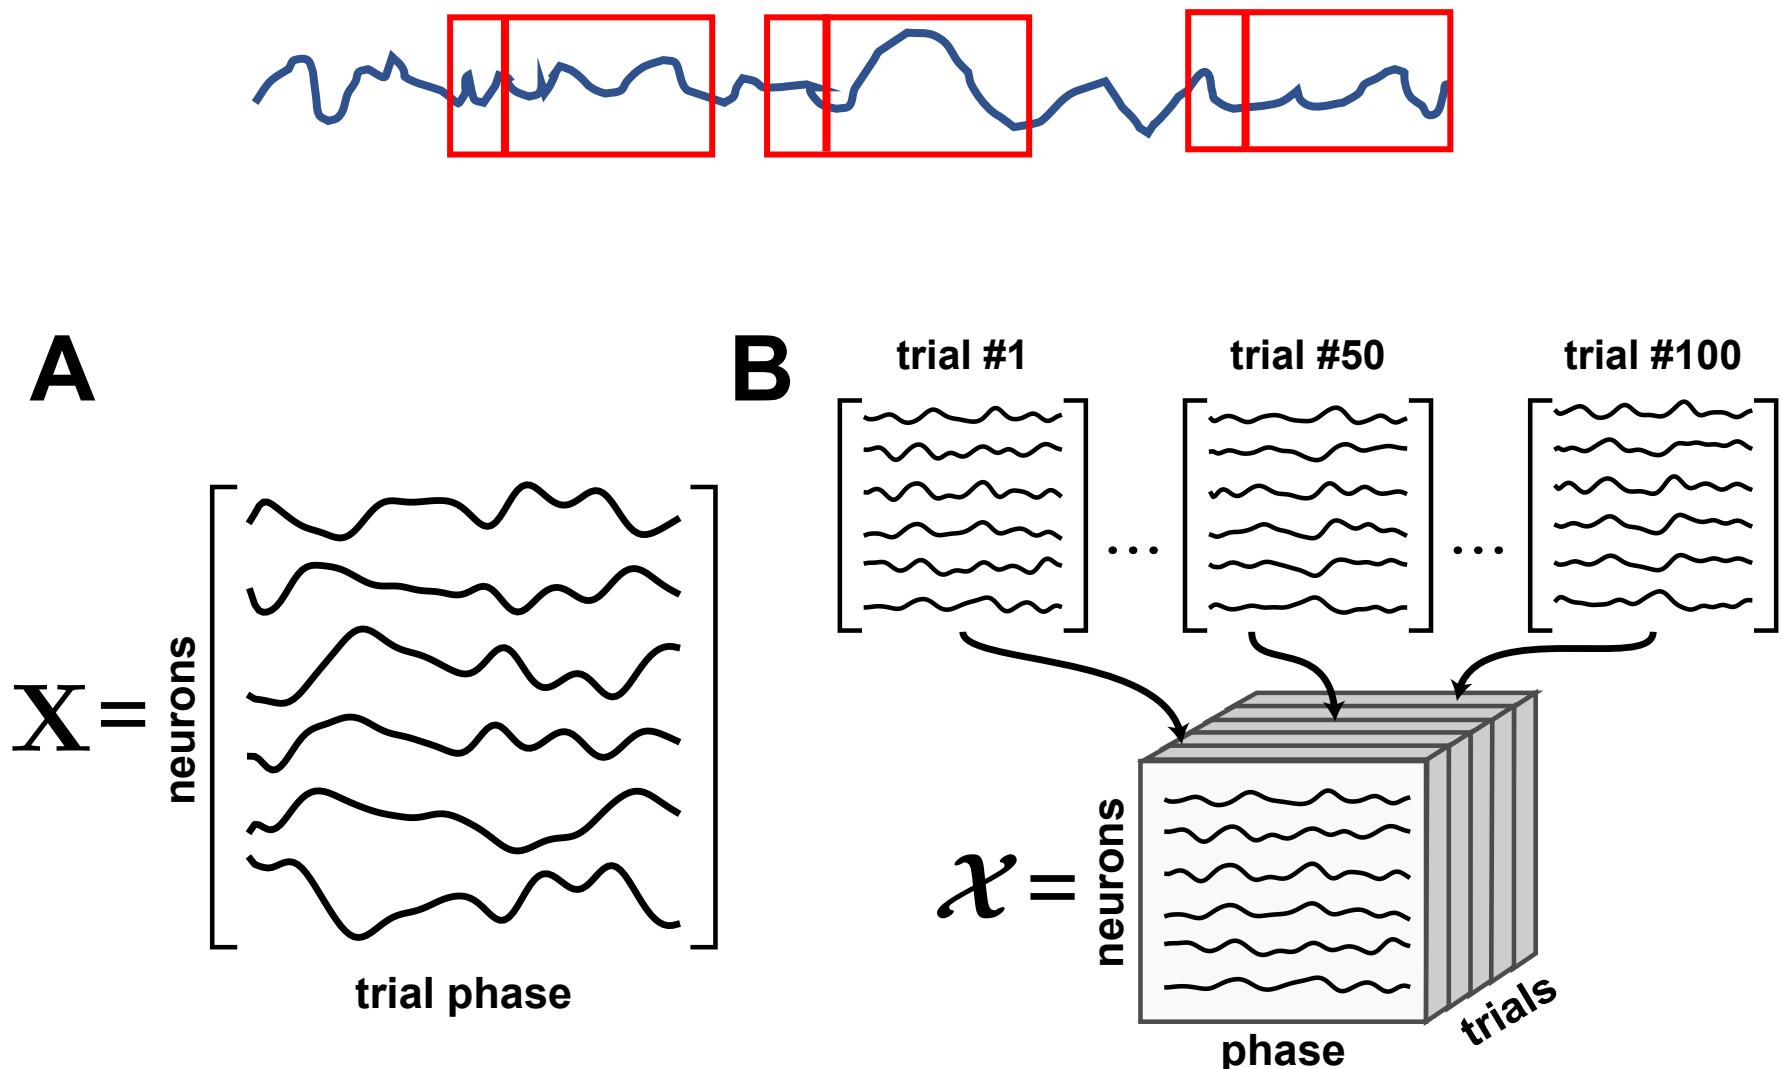
Paiva et al., Neural Computing and Applications, 2010

Kernels:

'auto', 'rectangular', 'gaussian', 'exponential', 'alpha'

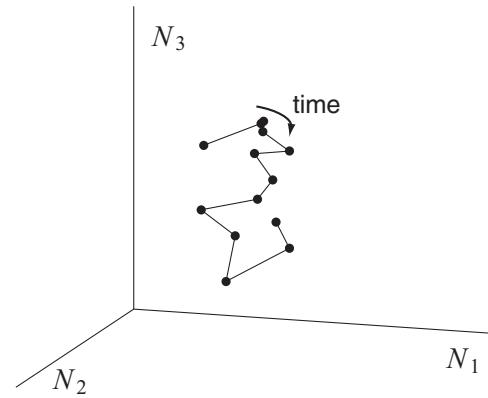
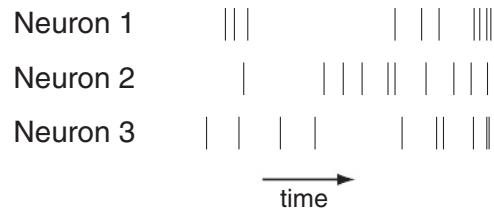
Auto: Gaussian with optimized standard deviation

H. Shimazaki and S. Shinomoto, "Kernel Bandwidth Optimization in Spike Rate Estimation," in Journal of Computational Neuroscience 29(1-2): 171–182, 2010



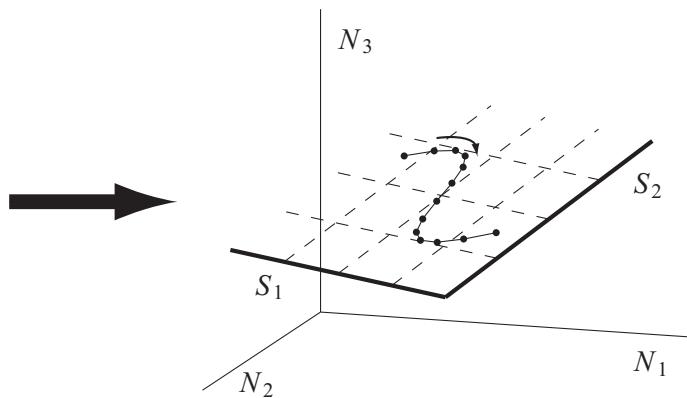
“tensor”

Standard Approach: Dimensionality reduction

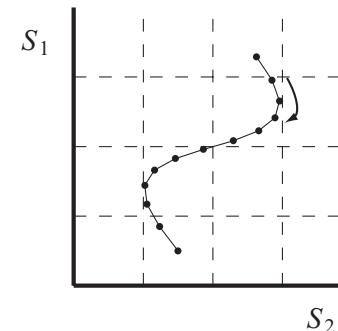


Spike trains

Noisy time series

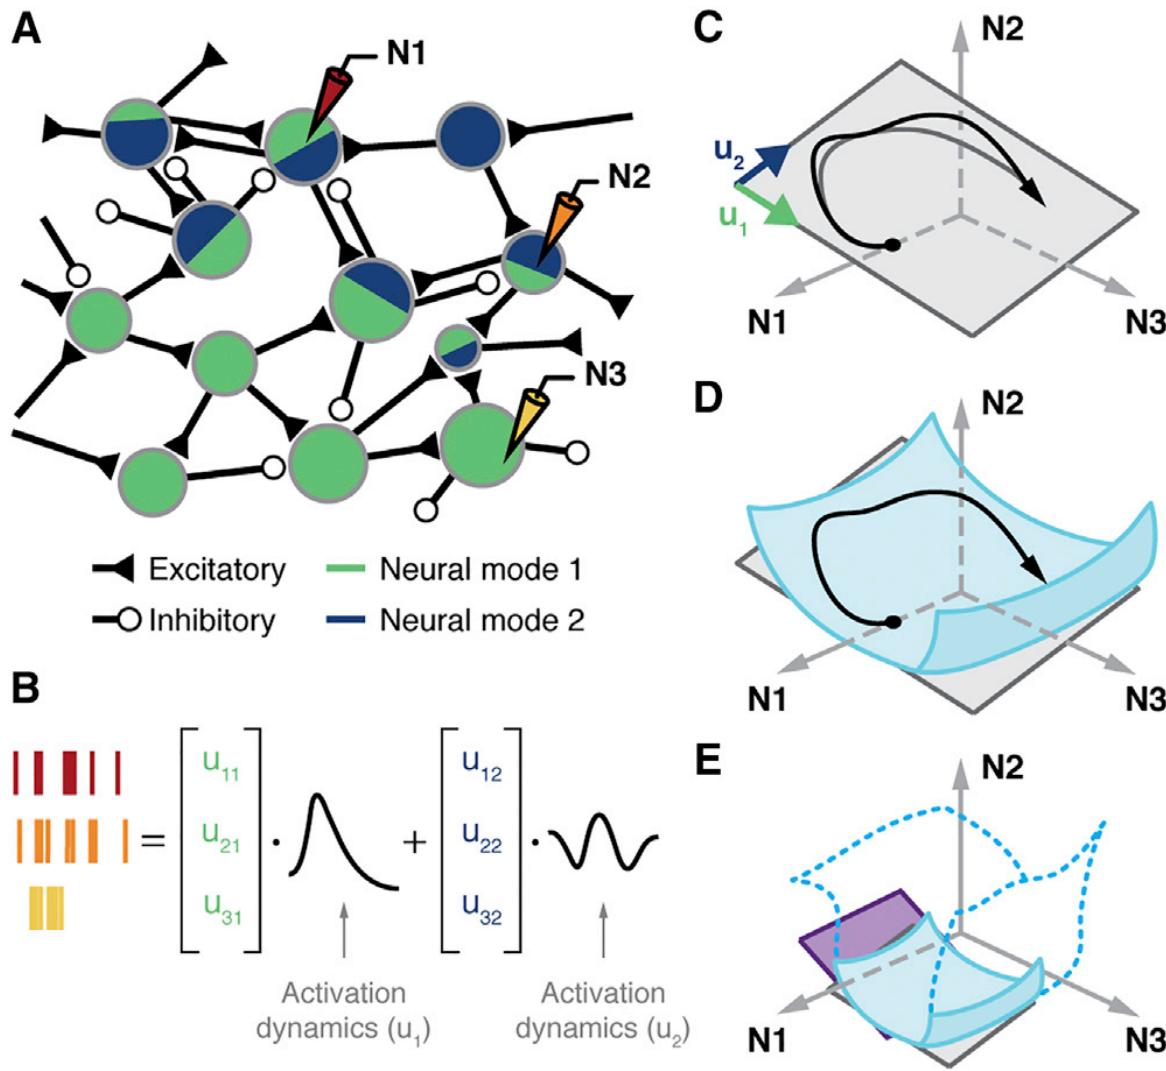


Denoised time series



Low-dimensional
time series

Standard Approach: Dimensionality reduction



Gallego et al., Neuron 2017

Standard Approach:

- State Estimation (Kalman Filter)
 - Uncertain Information → Educated Guess about What System Will do Next
 - Example:
 - Robot moving in one direction with **fairly** constant speed
 - It has a GPS which tells us the exact position with 10 meters **uncertainty**
 - We define its state as $\vec{x}_k = (\vec{p}, \vec{v})$
 - Kinematics **OR** GPS ?
 - Using all available information would give us a better answer than either estimation



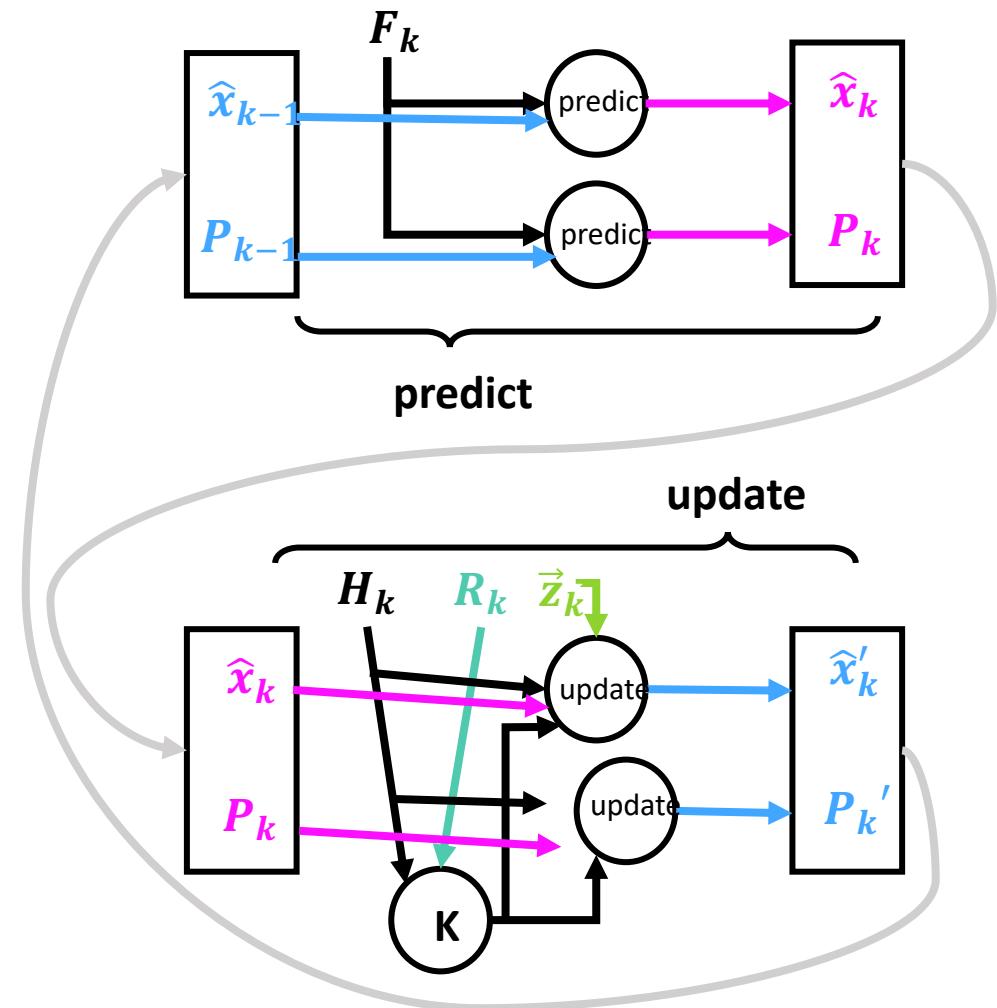
State Vector (x) comprises

- Position
- Velocity
- Acceleration

Predict next step (position =
velocity * time_delta)

Update with neural input (z)

Optional: z can include history
with previous time bins (taps)



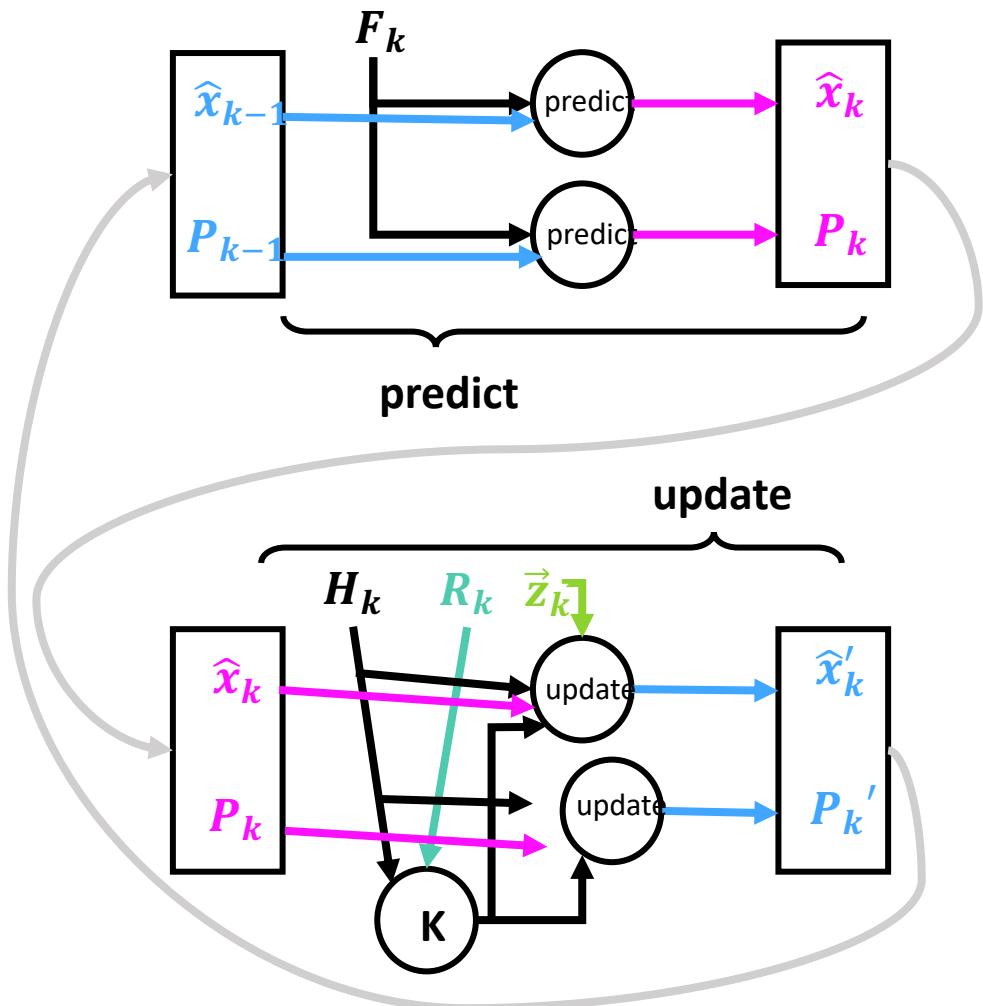
- State Estimation (Kalman Filter)

- Predict

- $\hat{x}_k = F_k \hat{x}_{k-1}$
- $P_k = F_k P_{k-1} F_k^T$

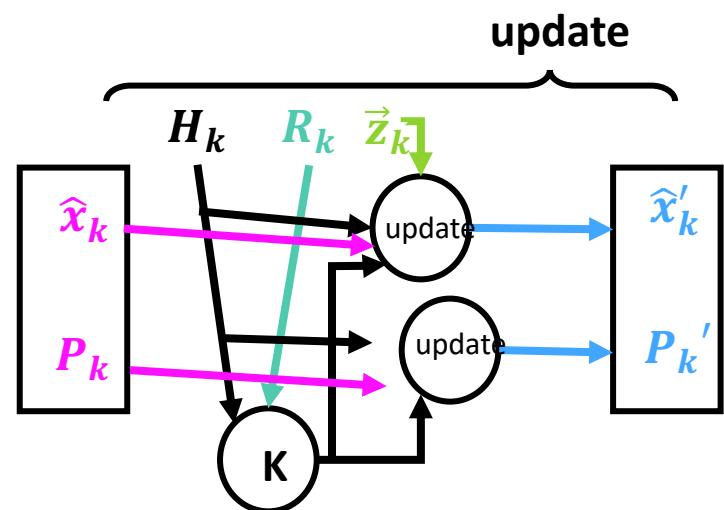
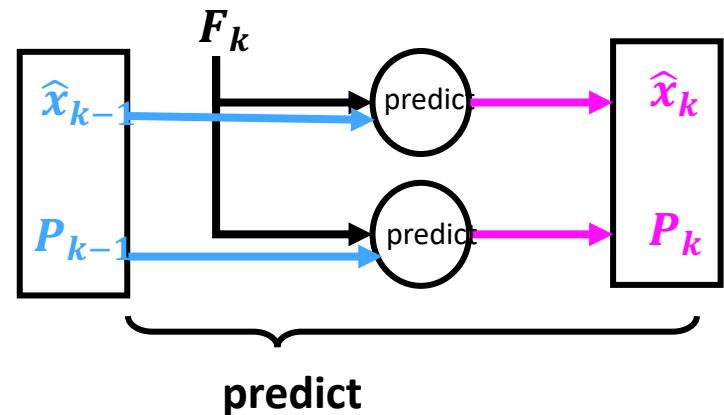
- Update

- $\hat{x}'_k = \hat{x}_k + K'(\vec{z}_k - H_k \hat{x}_k)$
- $P_k' = P_k - K' H_k P_k H_k^T$
- $K' = P_k H_k^T (H_k P_k H_k^T + R_k)^{-1}$



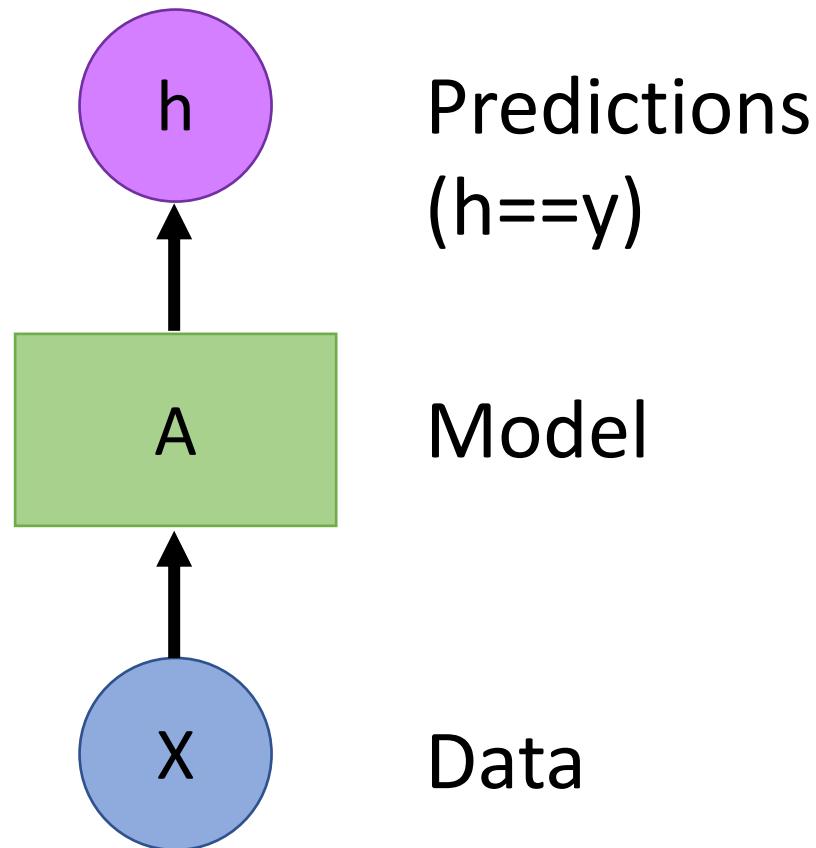
- State Estimation (Kalman Filter)

- Predict
 - New estimate is predicted from previous best estimate
 - New uncertainty is predicted from old uncertainty
- Update
 - New best estimate is predicted from new estimate and sensor input
 - Overall new uncertainty is predicted from new uncertainty and sensor noise

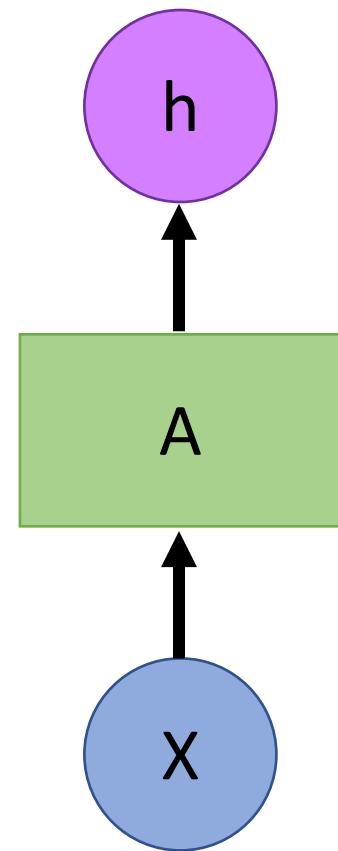
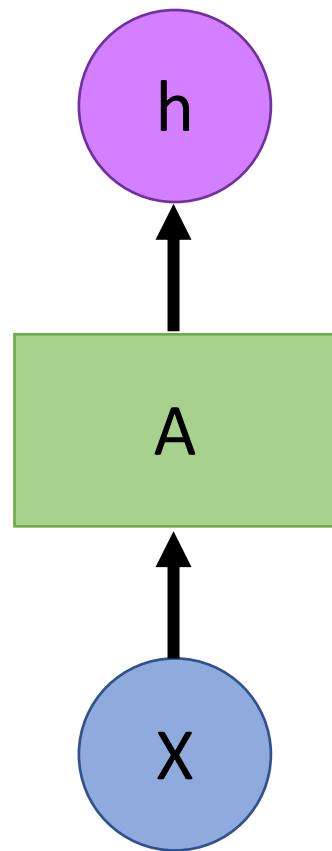
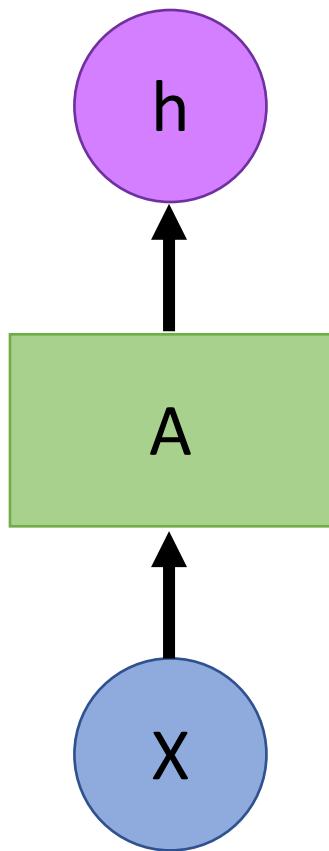


- ❖ F_k is how new prediction is related to previous (i.e. Kinematics)
- ❖ H_k is how we measure our state from sensor inputs

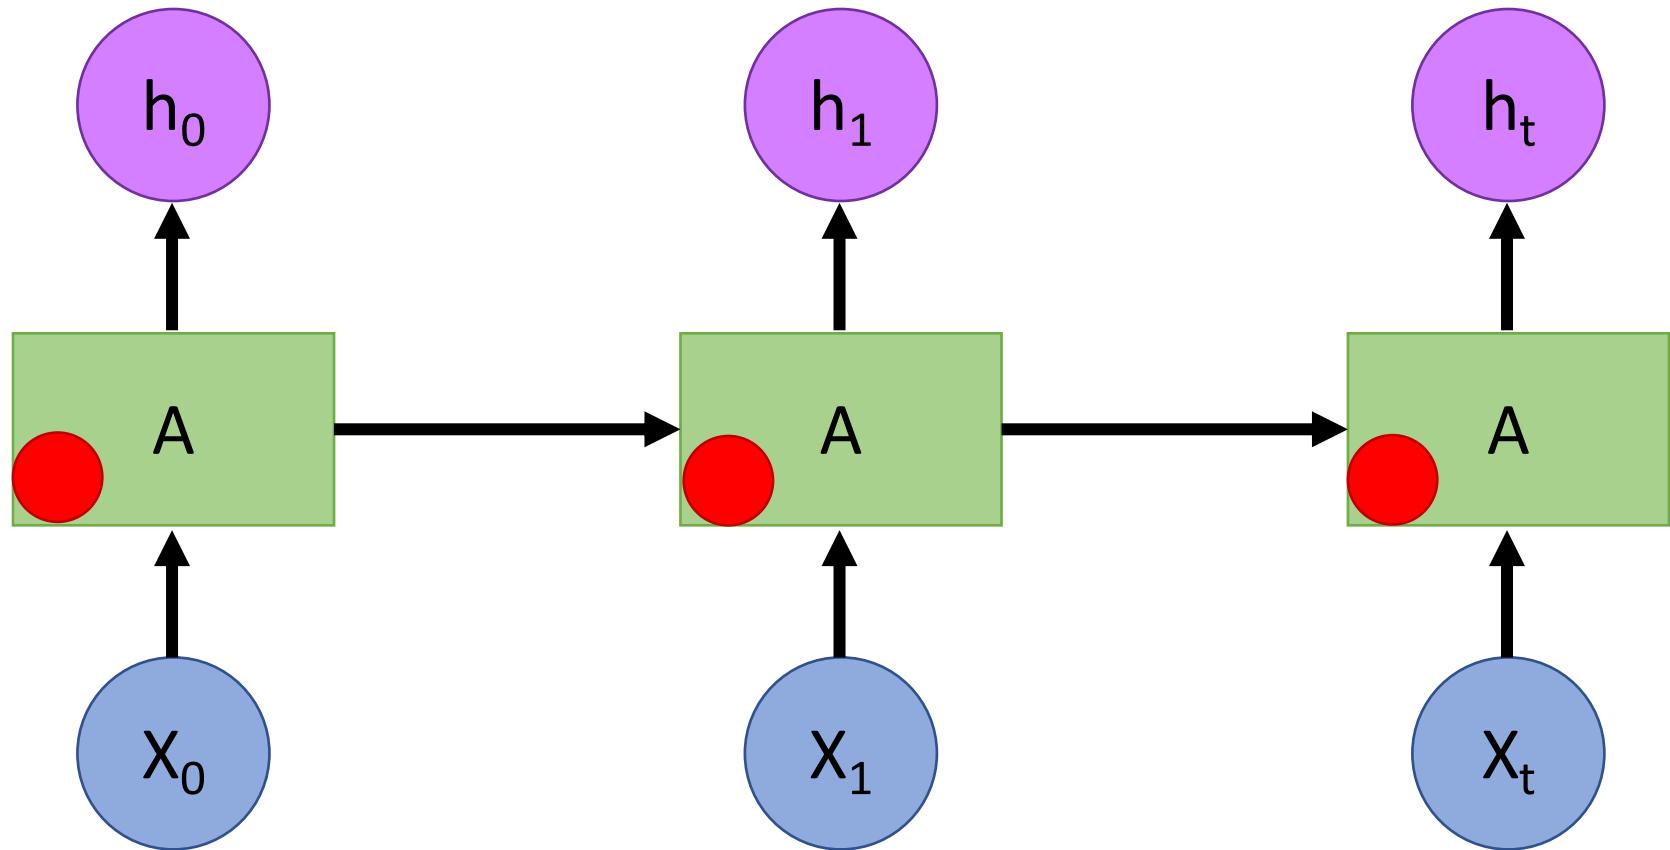
Vanilla CNN is stateless: no memory



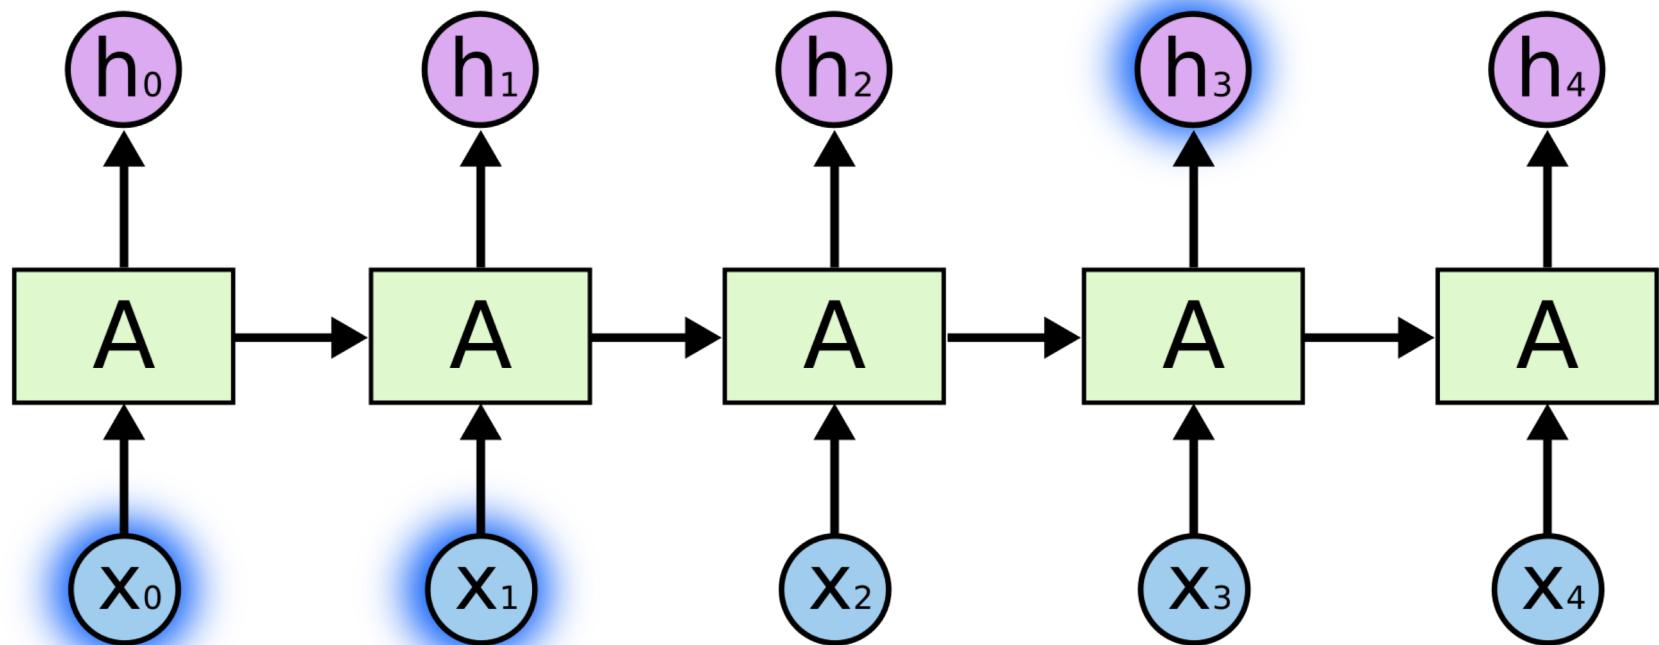
Vanilla CNN is stateless: no memory



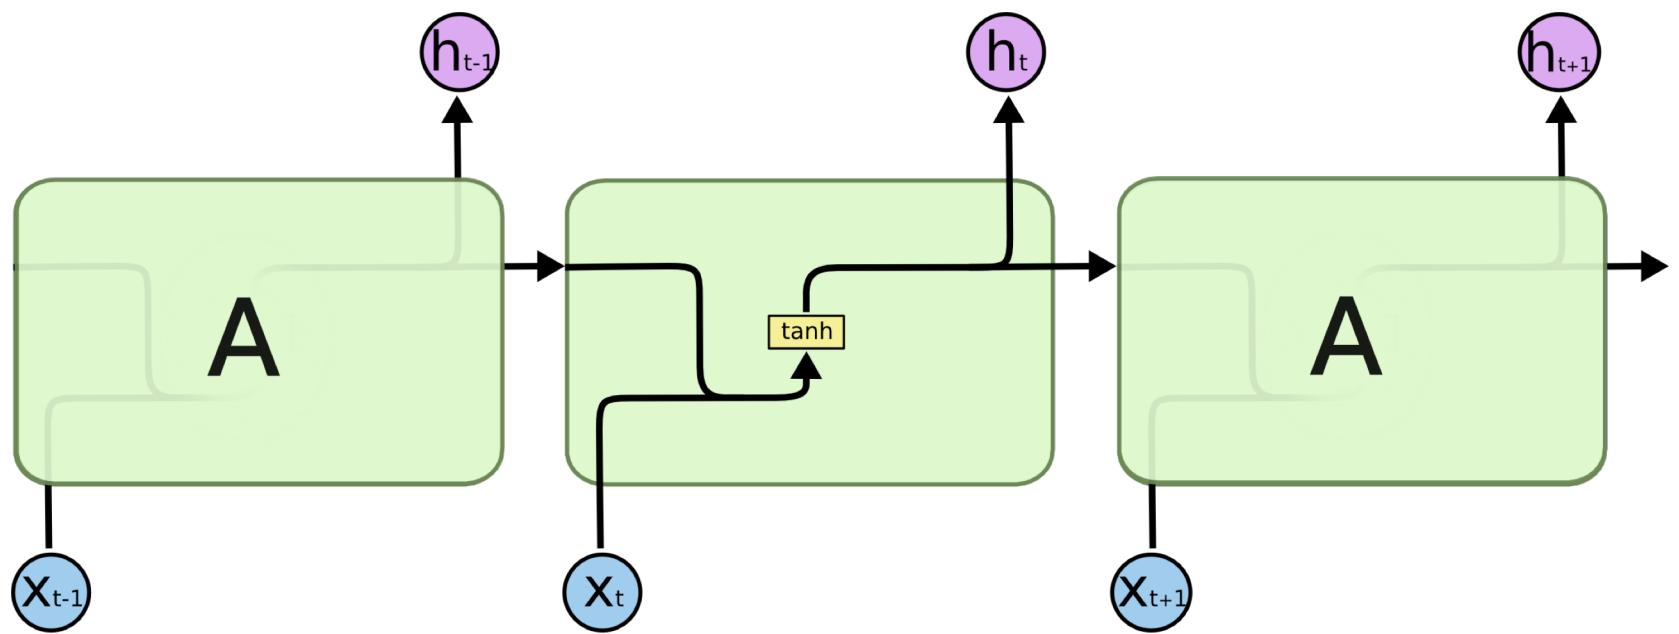
A model that remembers its (recent) data history is desirable.



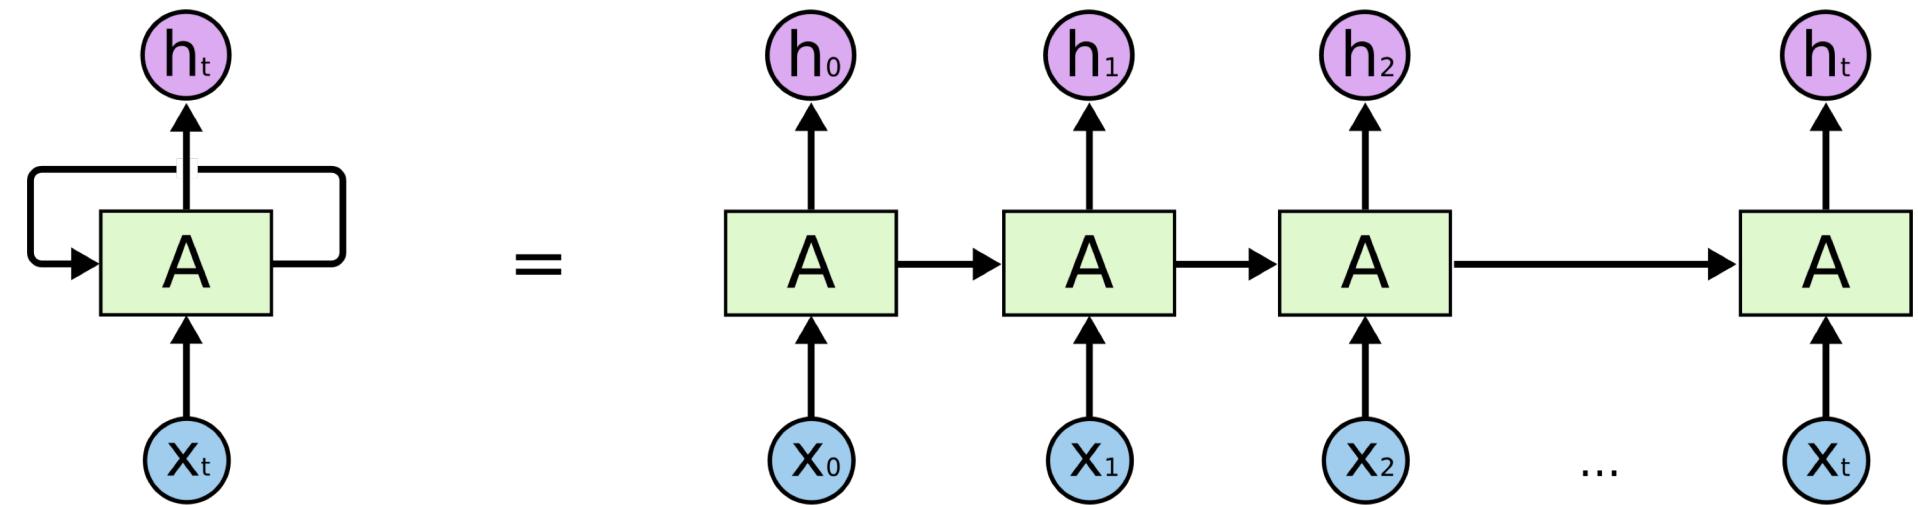
Recurrent neural networks (RNNs)
can use past information to make the
current prediction



[Figure source](#)



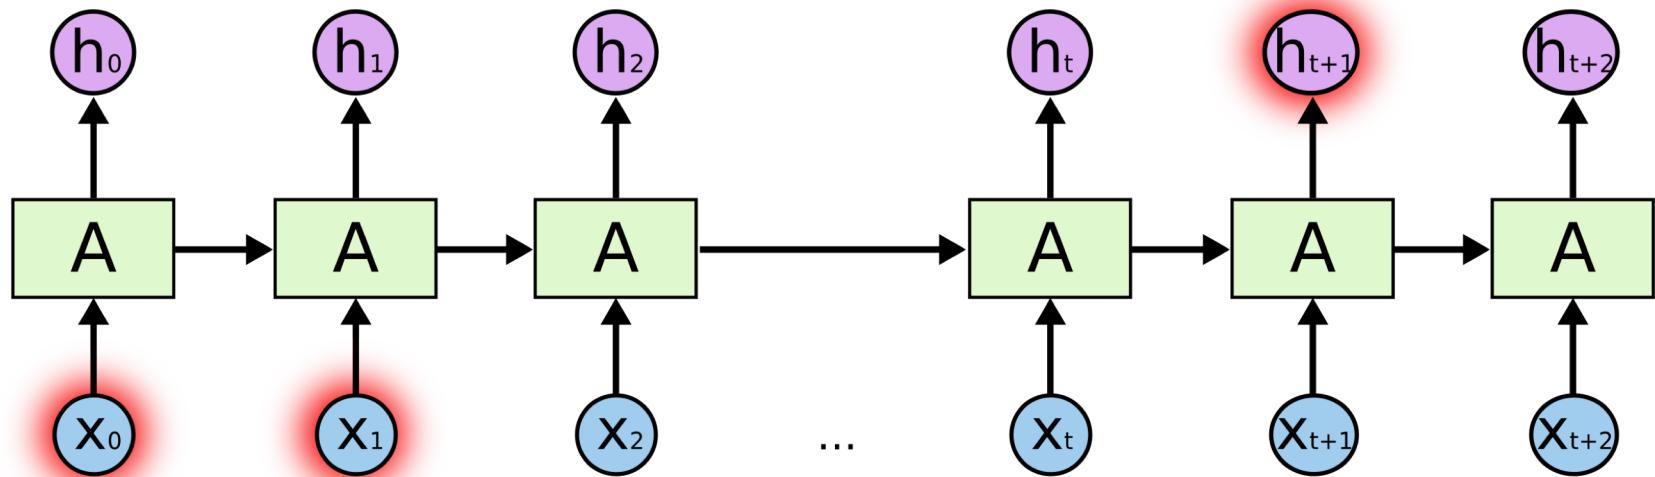
A model that maintains state and loops over its input is represented with a single looping model



- 05_01_intro_to_RNNs.ipynb
- When that's done, open 05_02 and start downloading data, then resume here.

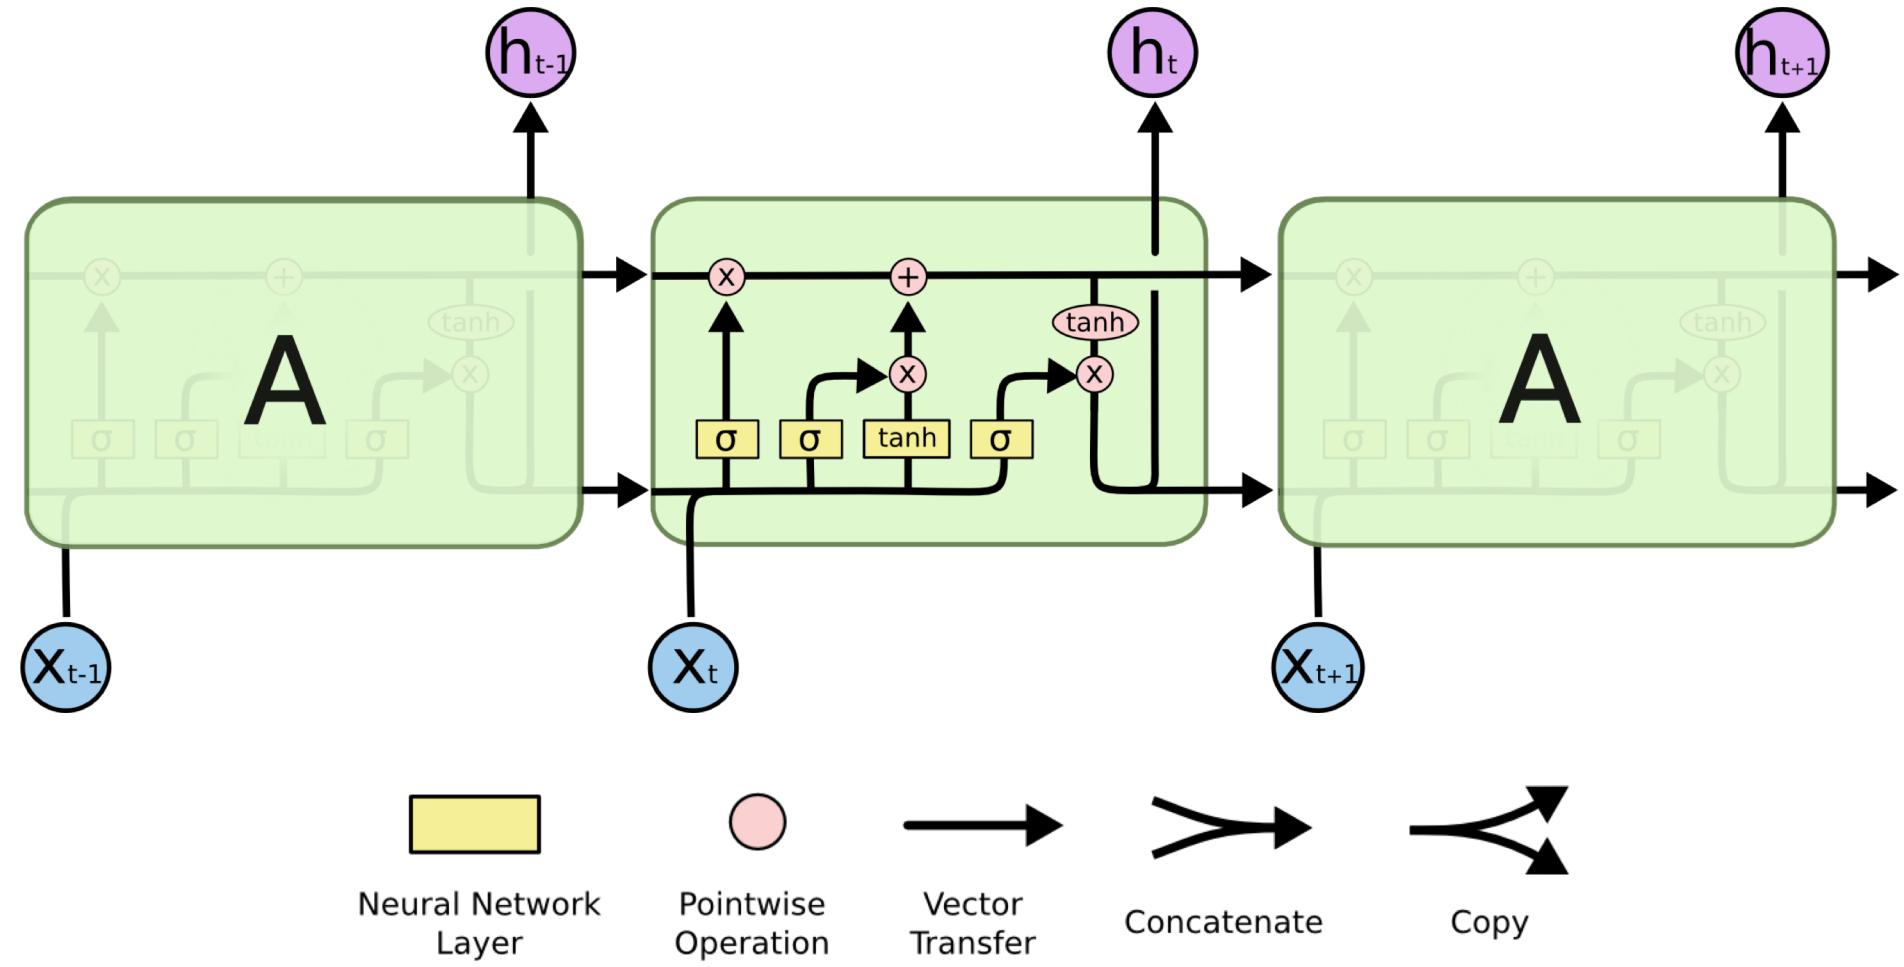
Learning Long-Term Dependencies with Gradient Descent is Difficult

Bengio et al., 1994

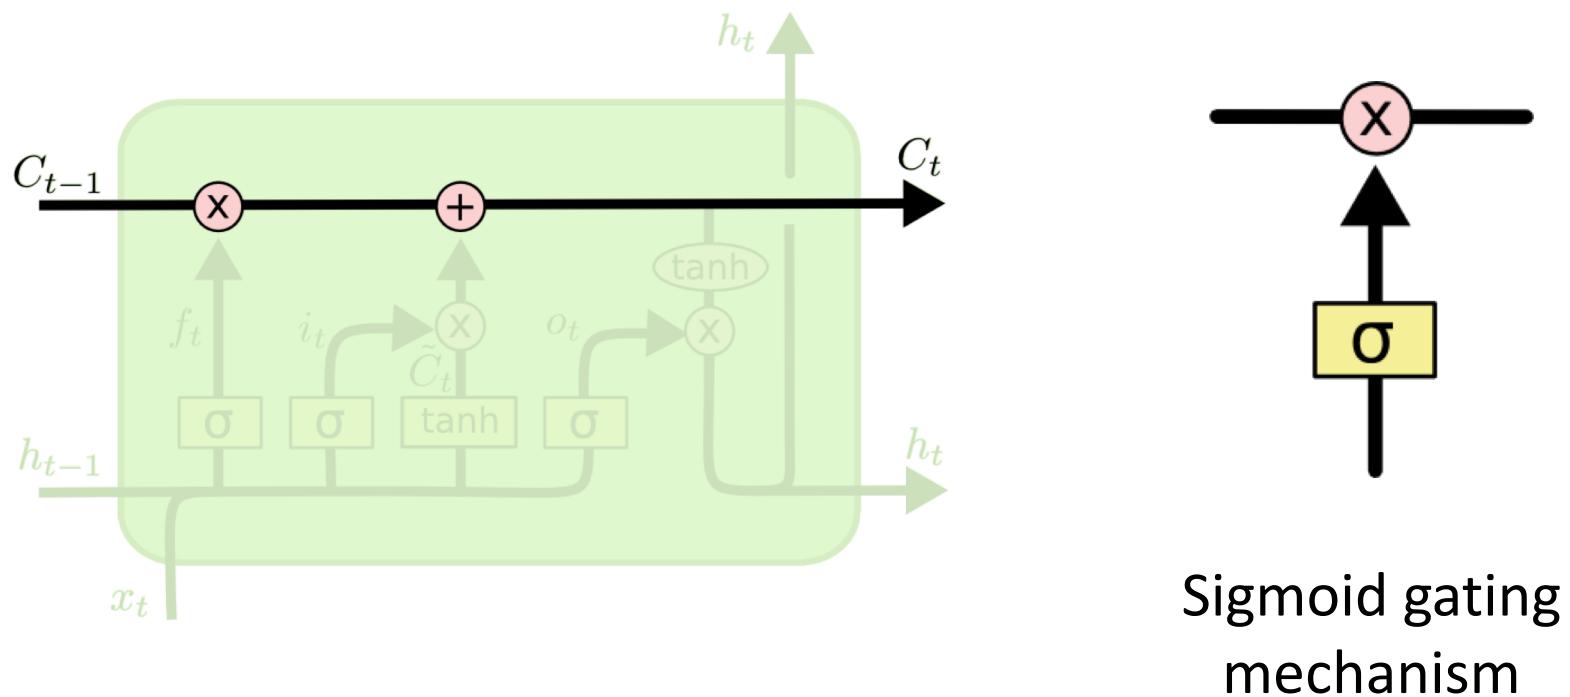


- In practice, simple RNNs cannot maintain information for more than a few steps
 - Exploding/Vanishing gradient

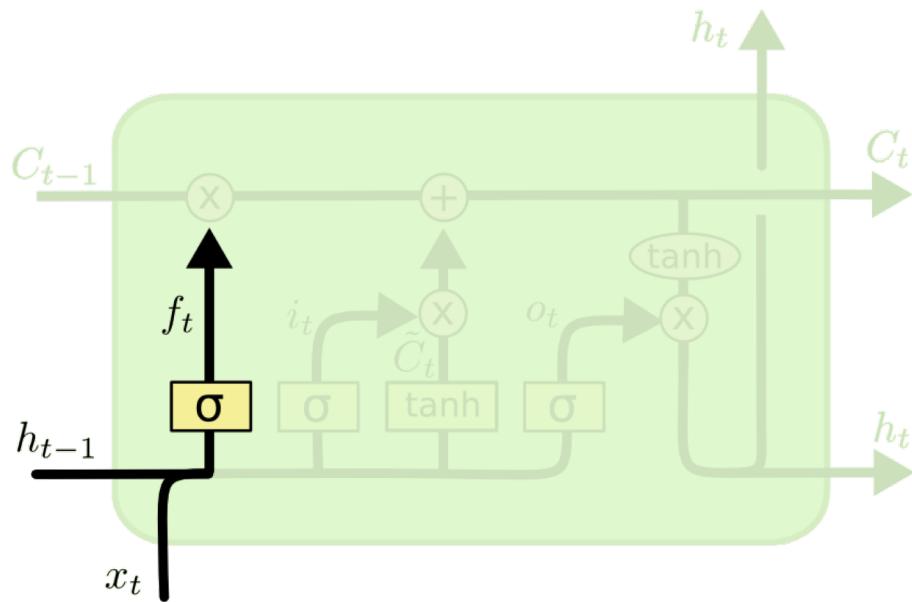
Long-Short Term Memory (LSTM) are a special kind of RNN capable of learning long(er)-term dependencies



The cell state C persists across iterations. Its contents may be modified, determined by gates.

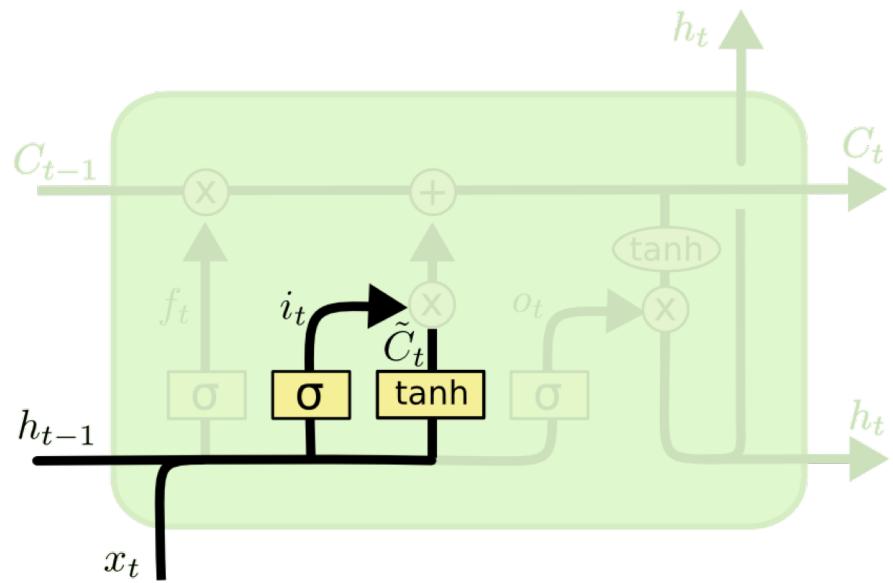


LSTM Step 1: Forget Gate



$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

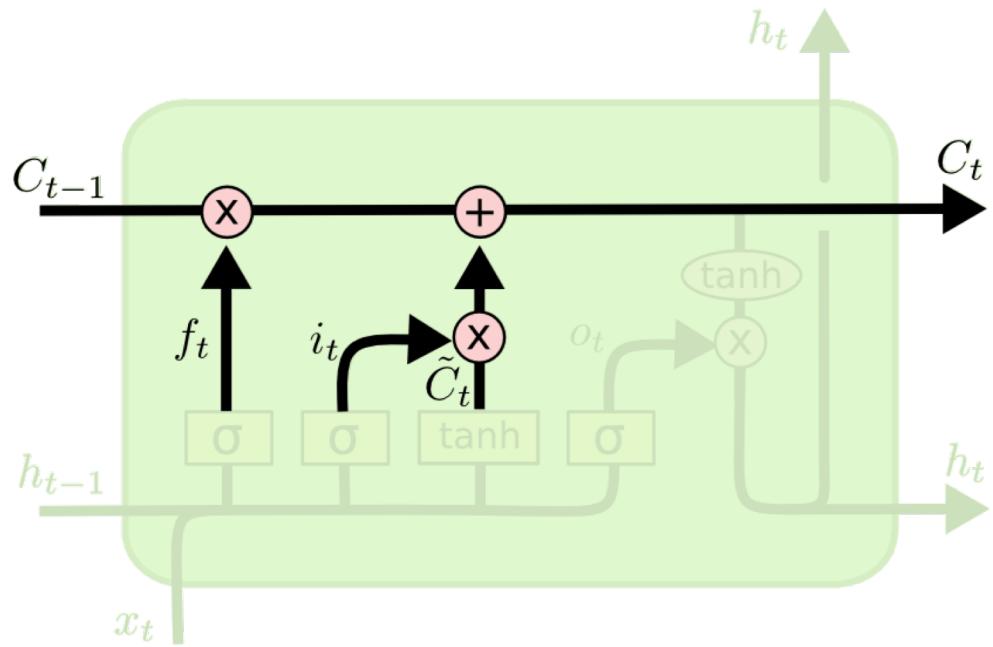
LSTM Step 2: Input Gate



$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

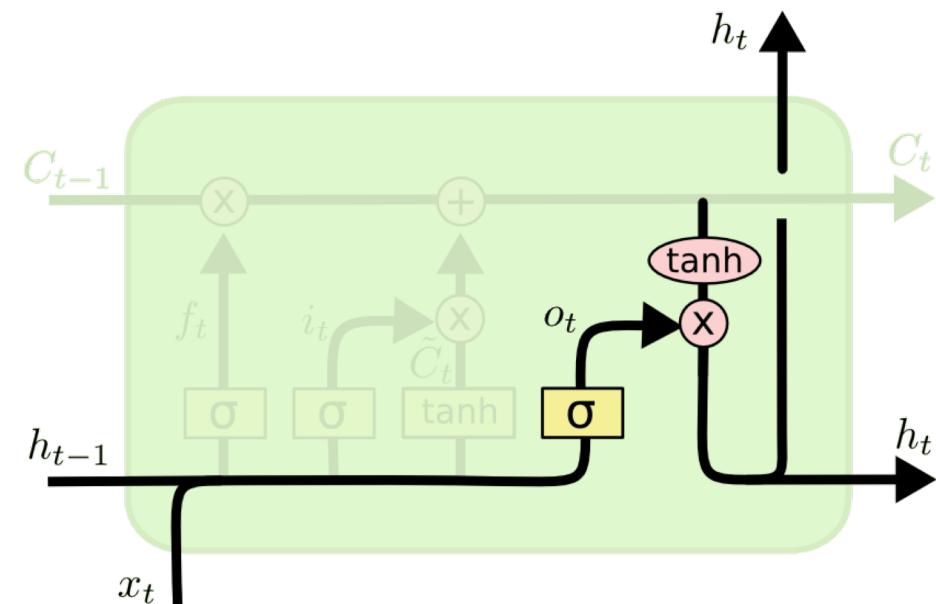
$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

LSTM Step 3: Modify State



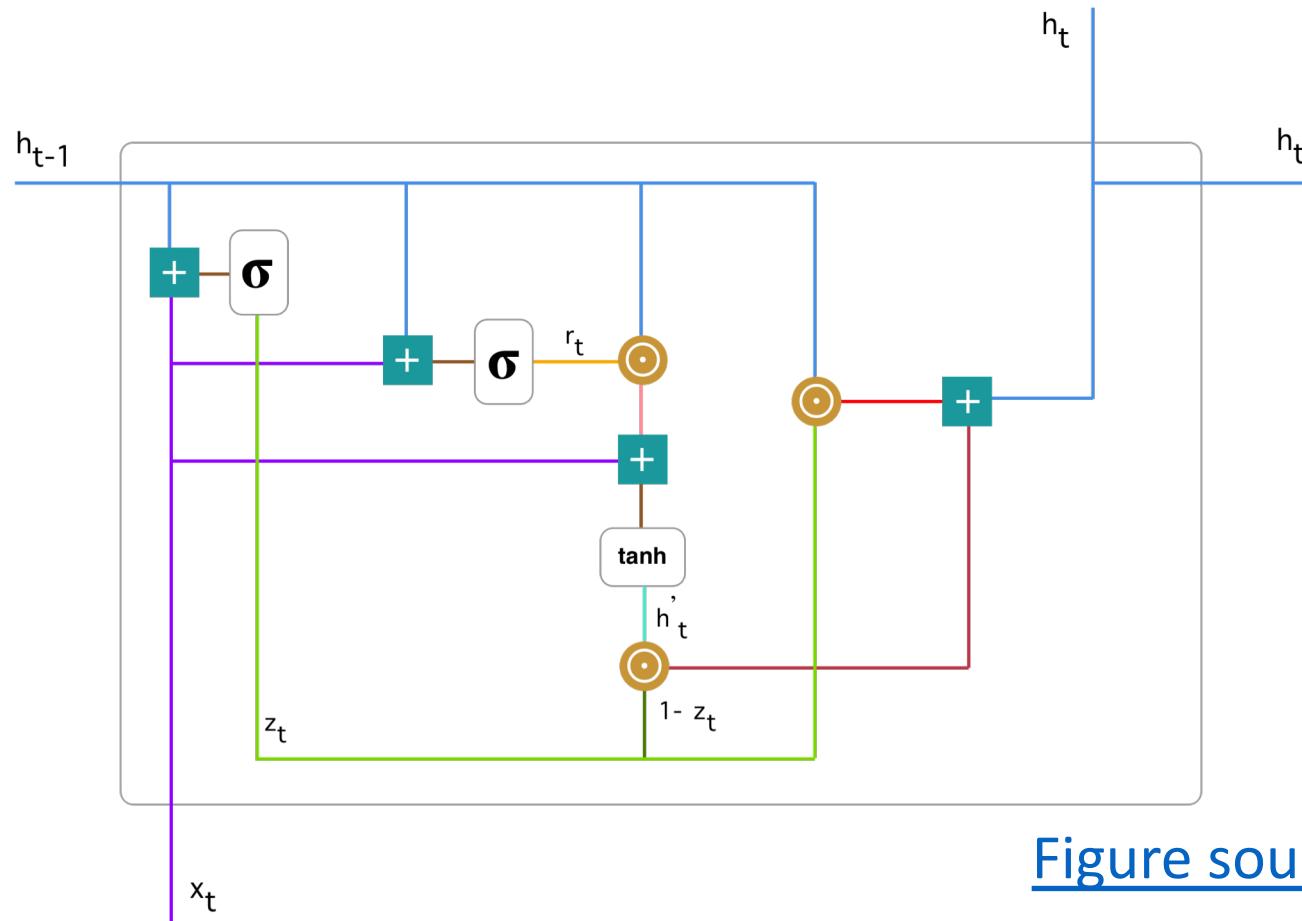
$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

LSTM Step 4: Combine (new) state and input to produce output



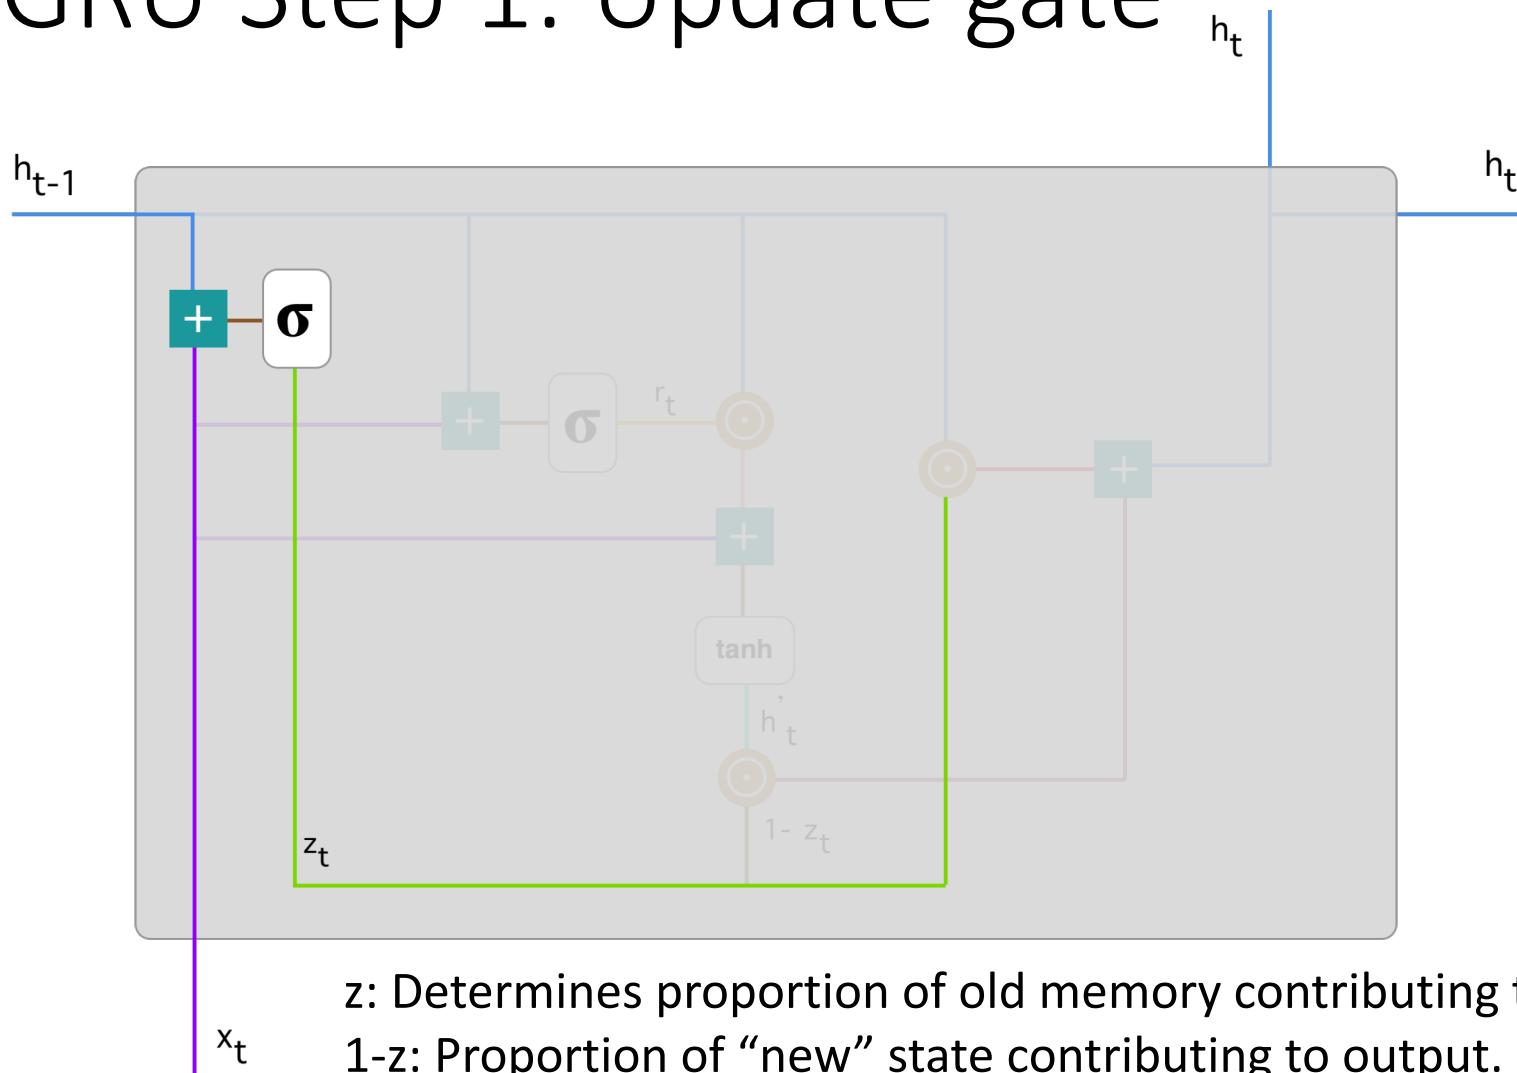
$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$
$$h_t = o_t * \tanh (C_t)$$

Gated Recurrent Unit (GRU)



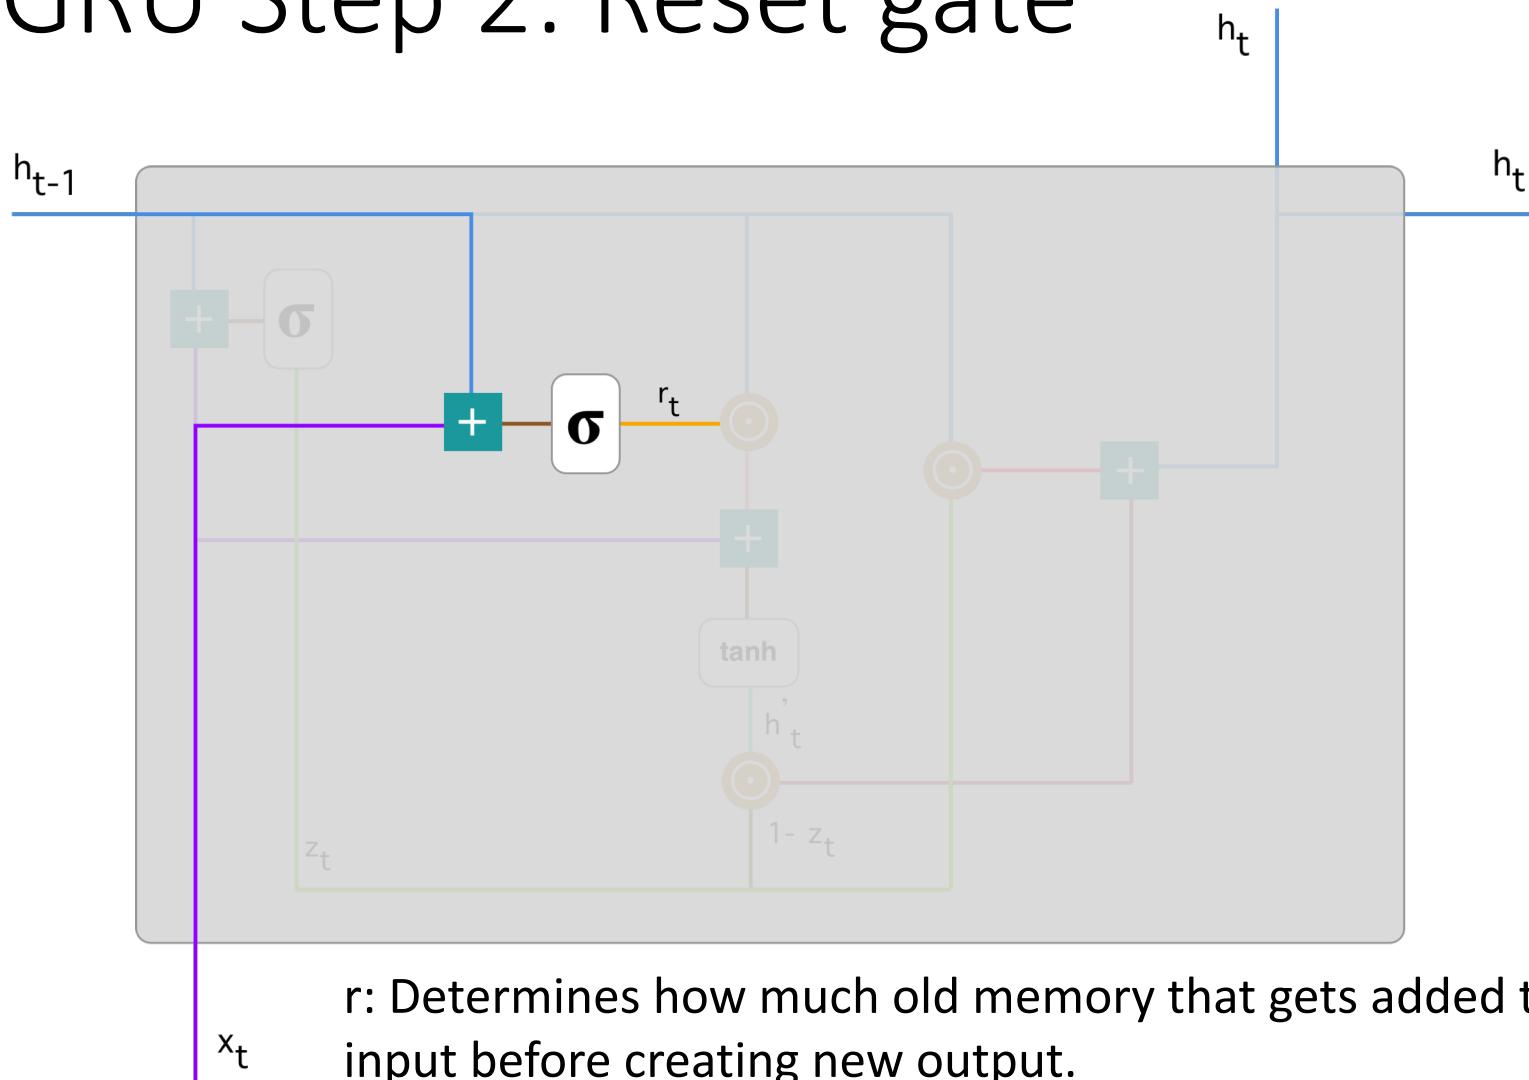
[Figure source](#)

GRU Step 1: Update gate



$$z_t = \sigma(W^{(z)}x_t + U^{(z)}h_{t-1})$$

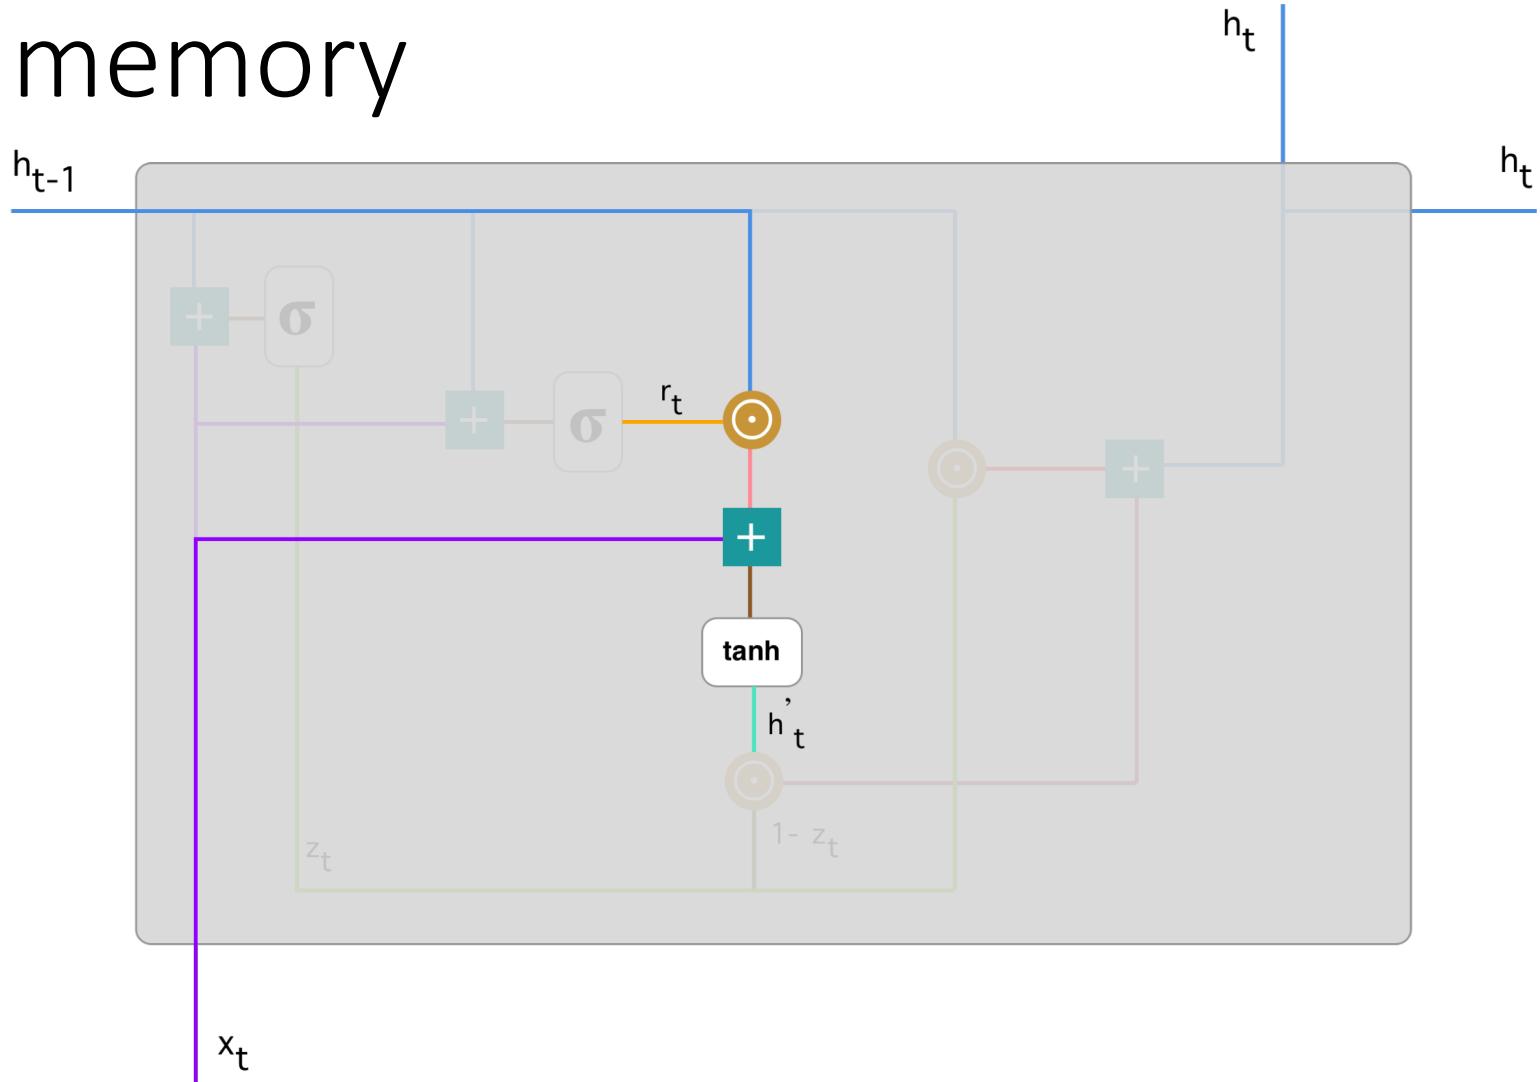
GRU Step 2: Reset gate



r_t : Determines how much old memory that gets added to new input before creating new output.

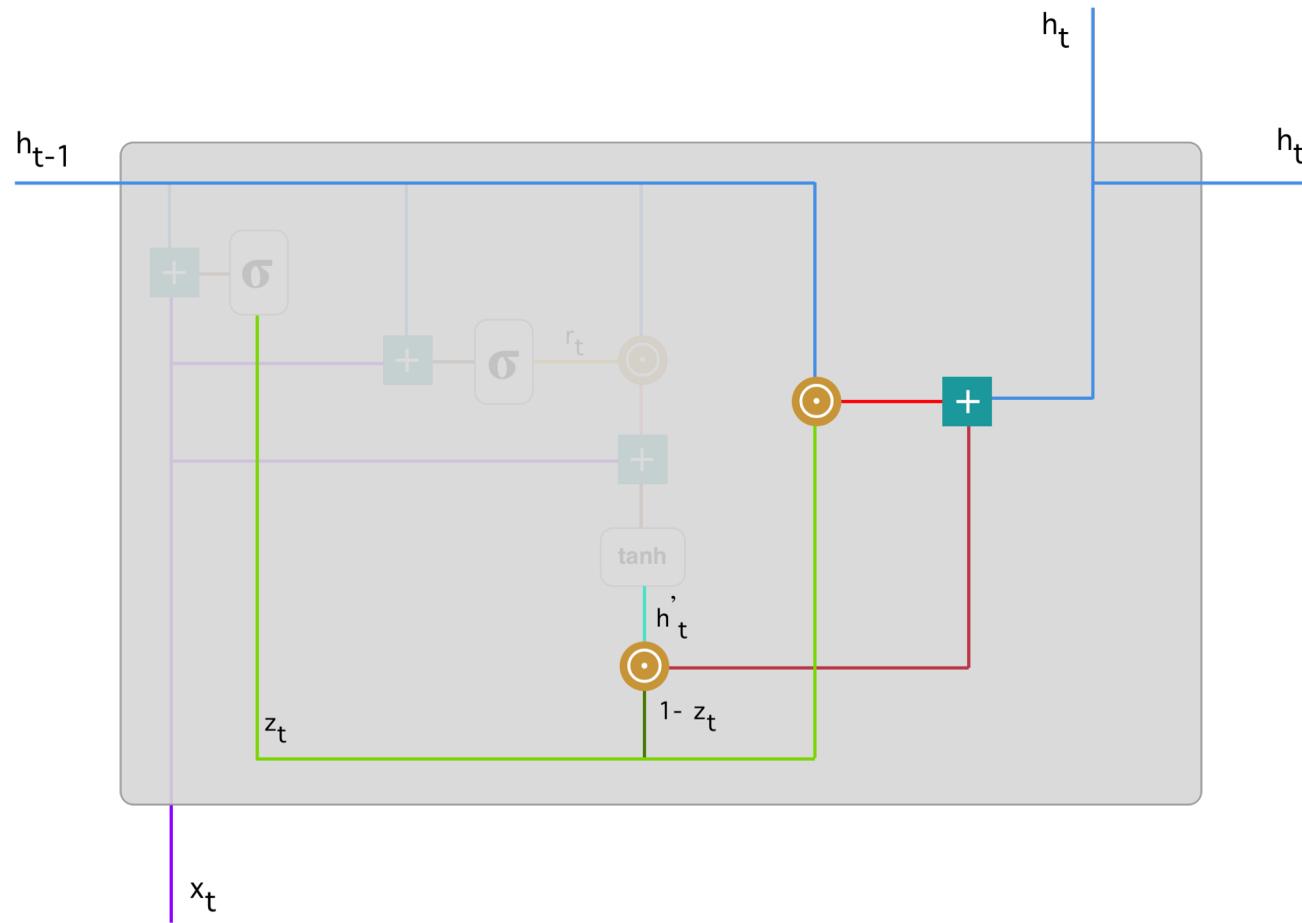
$$r_t = \sigma(W^{(r)}x_t + U^{(r)}h_{t-1})$$

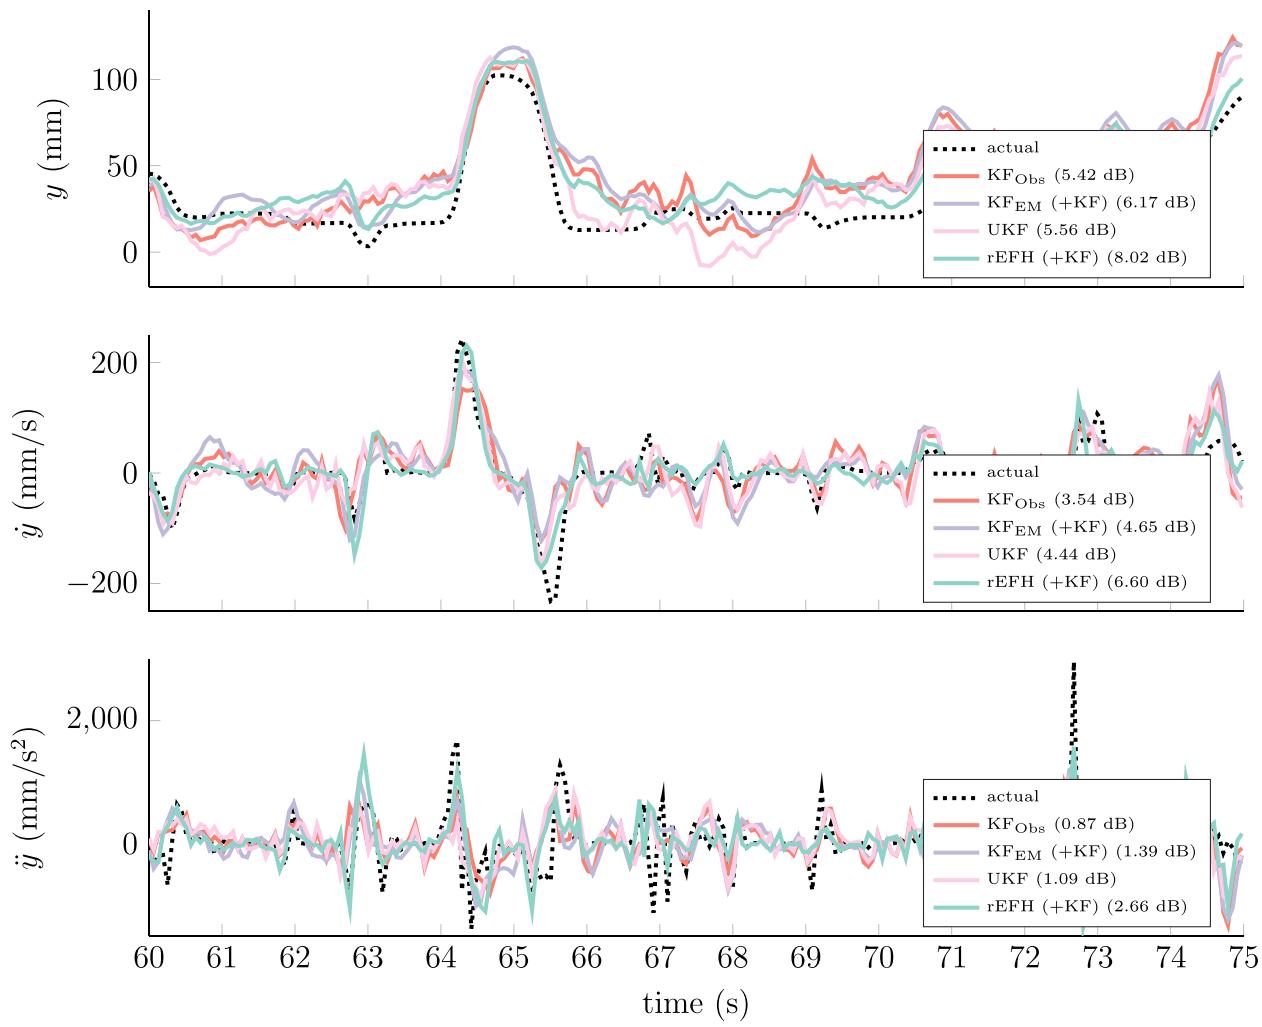
GRU Step 3: Combine input and memory



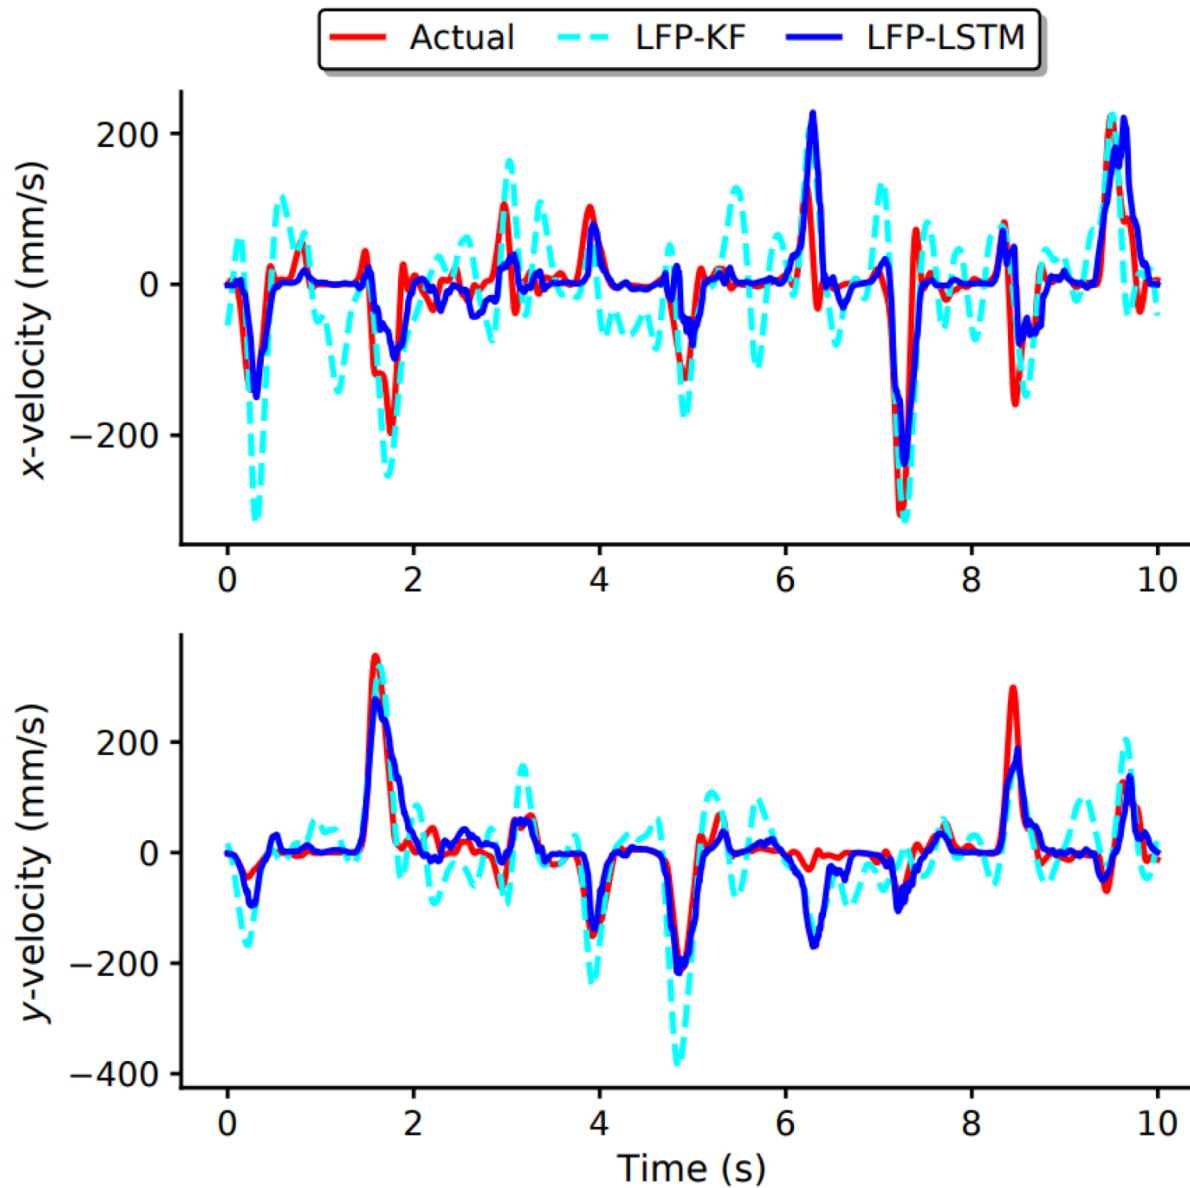
$$h'_t = \tanh(Wx_t + r_t \odot Uh_{t-1})$$

GRU Step 4: Update memory





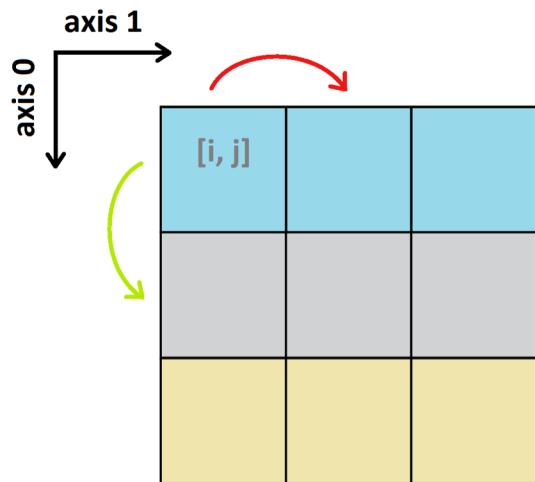
Result from Makin et al using same dataset



Result from Ahmadi et al., arXiv 2019

- 05_02_LSTM_and_GRU

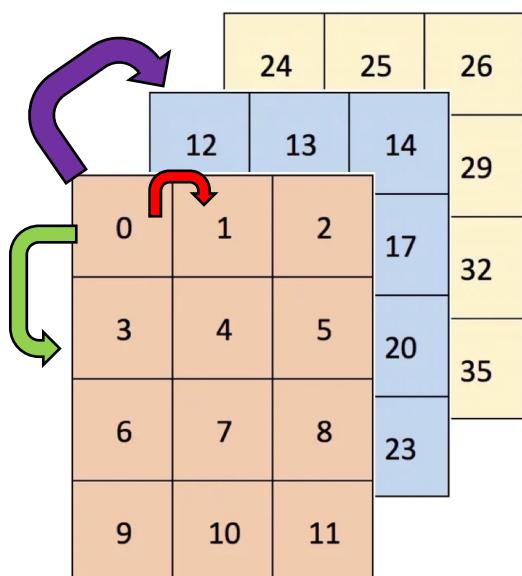
Array representation



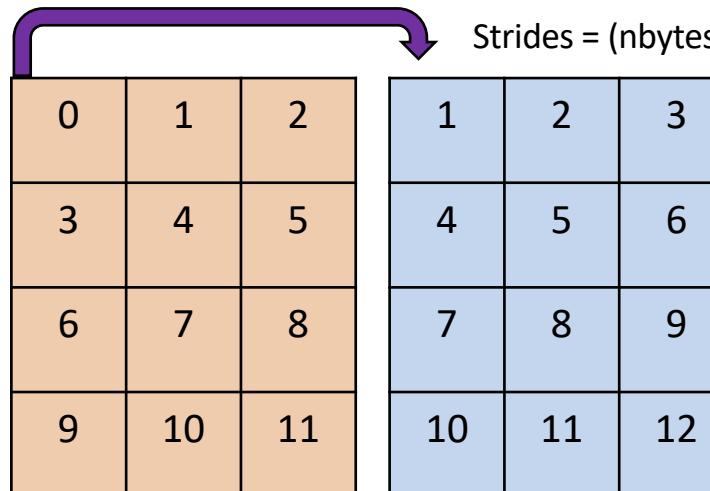
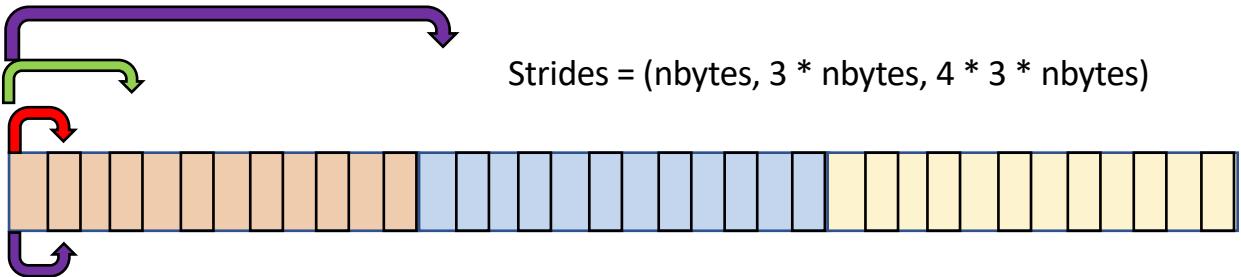
Contiguous block of memory



offset: $i * \text{strides}[0] + j * \text{strides}[1]$



Strides = (nbytes, 3 * nbytes, 4 * 3 * nbytes)



Strides = (nbytes, 3 * nbytes, nbytes)

Decoding Movements from Cortical Ensemble Activity Using a Long Short-Term Memory Recurrent Network

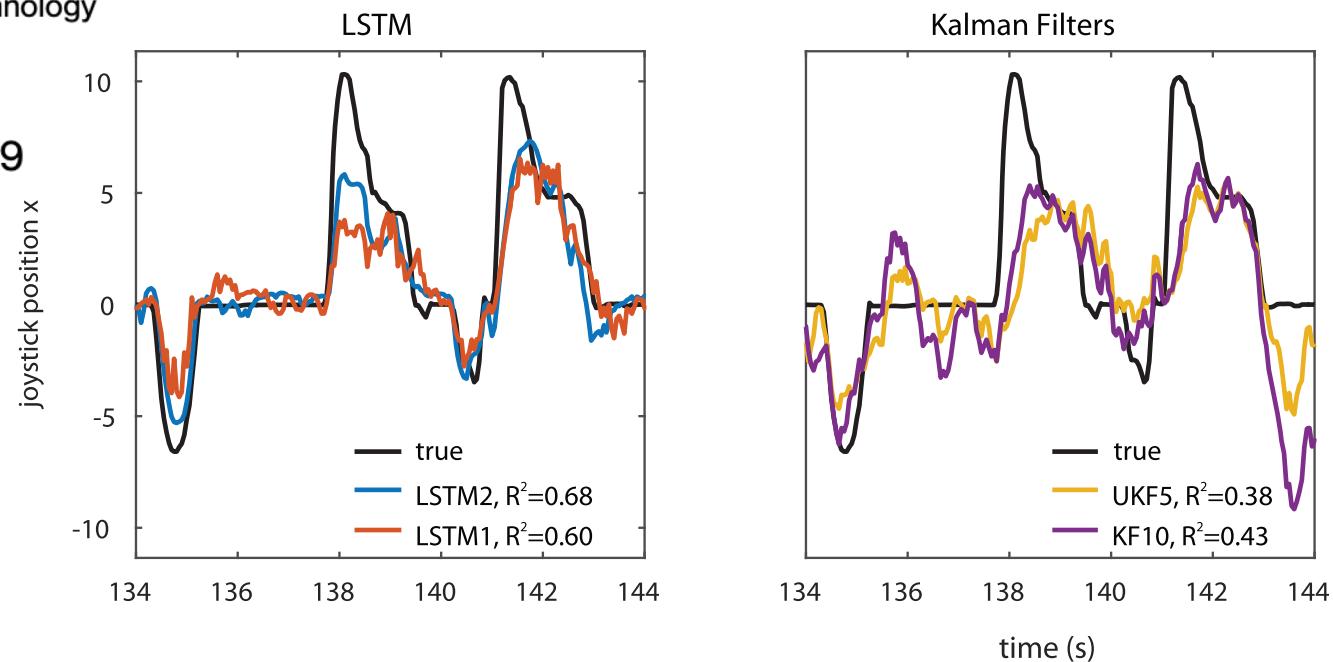
Po-He Tseng, Núria Armengol Urpi, Mikhail Lebedev and Miguel Nicolelis

Posted Online May 21, 2019

https://doi-org.proxy.bib.uottawa.ca/10.1162/neco_a_01189

(A)
© 2019 Massachusetts Institute of Technology

Neural Computation
Volume 31 | Issue 6 | June 2019
p.1085-1113



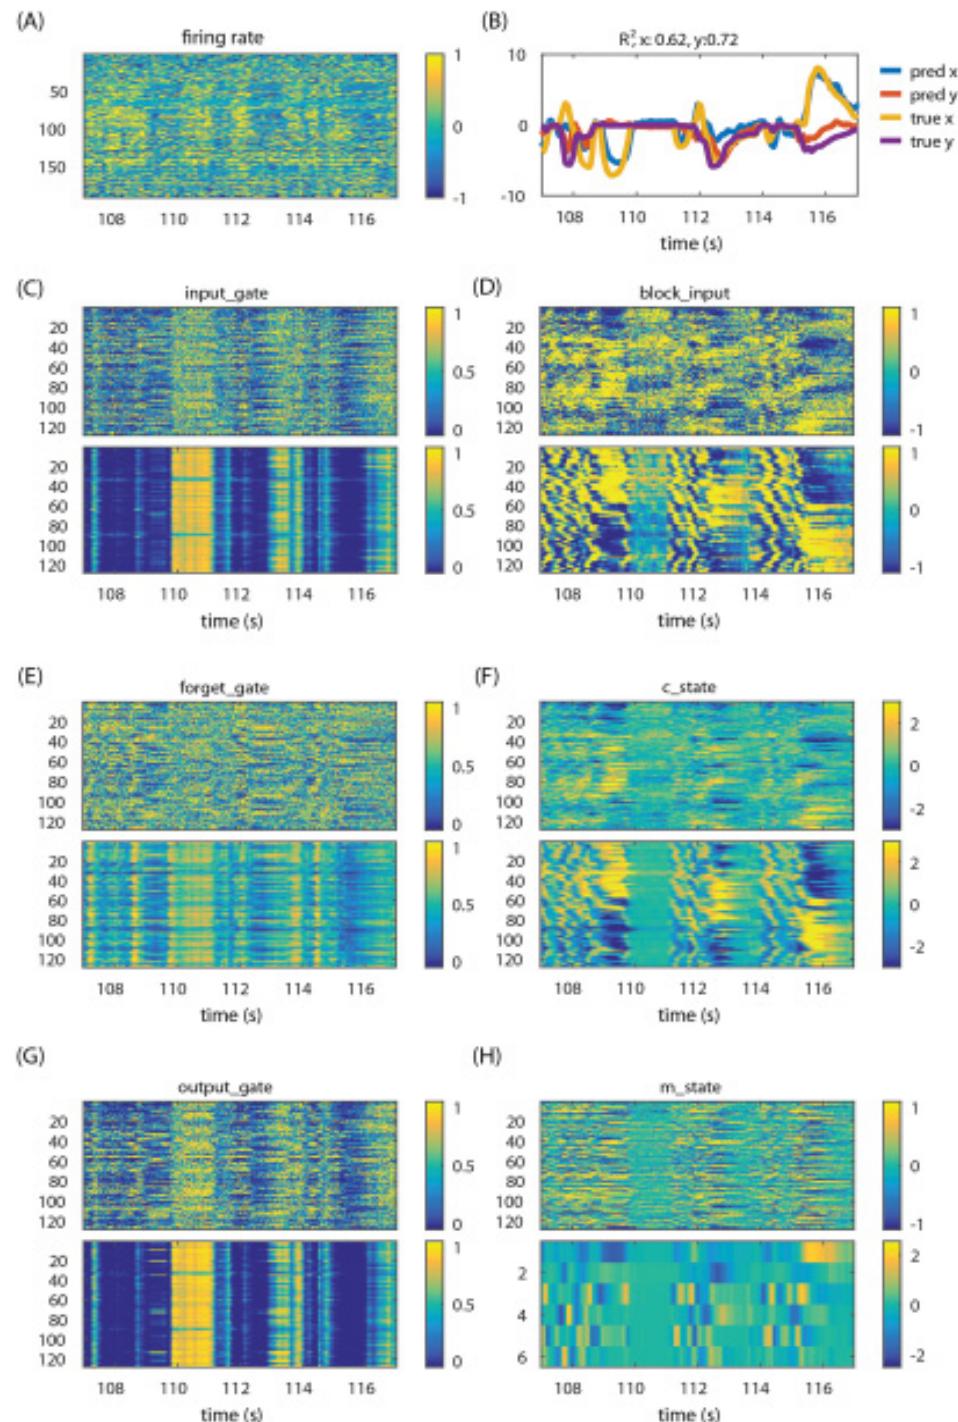
Model internal values in LSTM layers 1 and 2

Layers presented in pairs, with layer 1 on top.

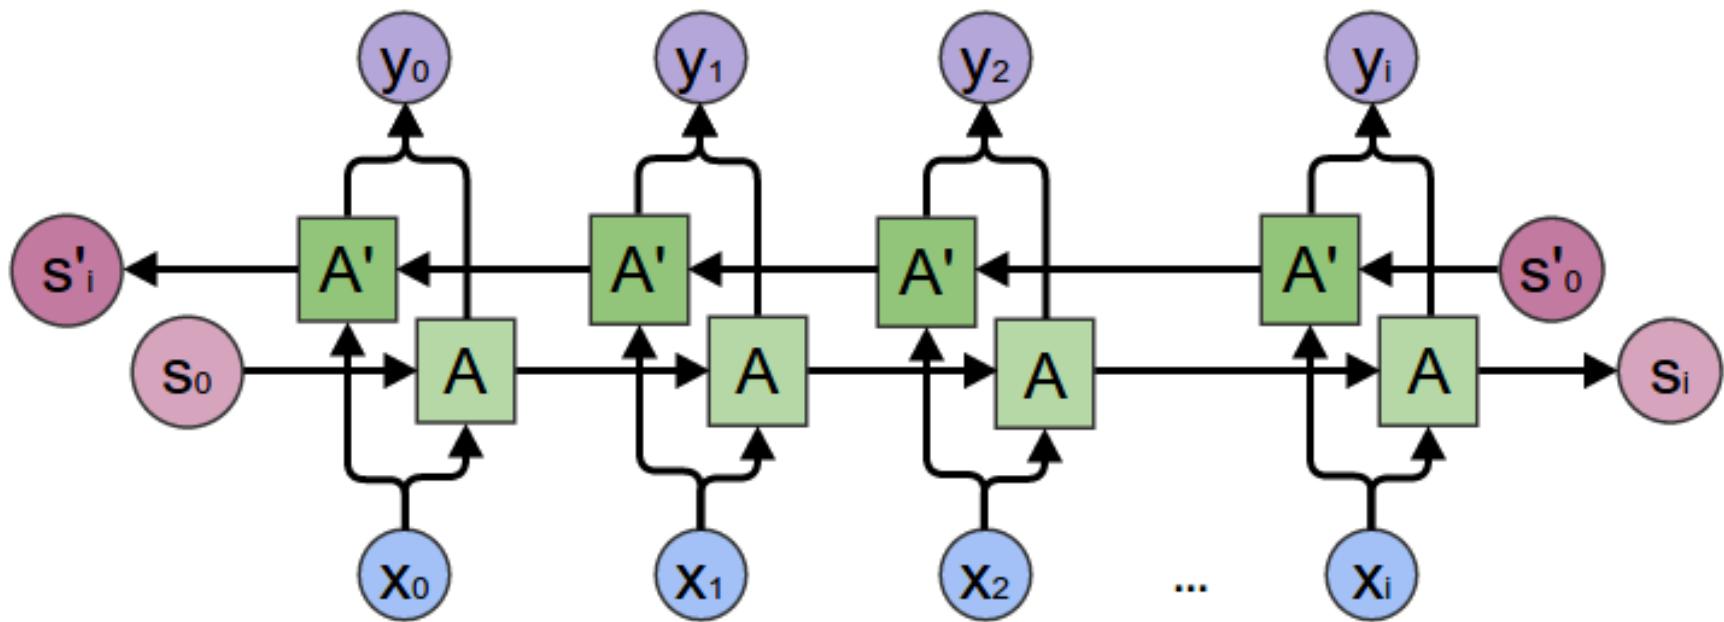
Layer 2 activations had higher contrast than layer 1, primarily because layer 2 gates were more fully closed.

(Not shown) layer 2 units were directionally tuned.

Does this tell us anything about how the brain works?

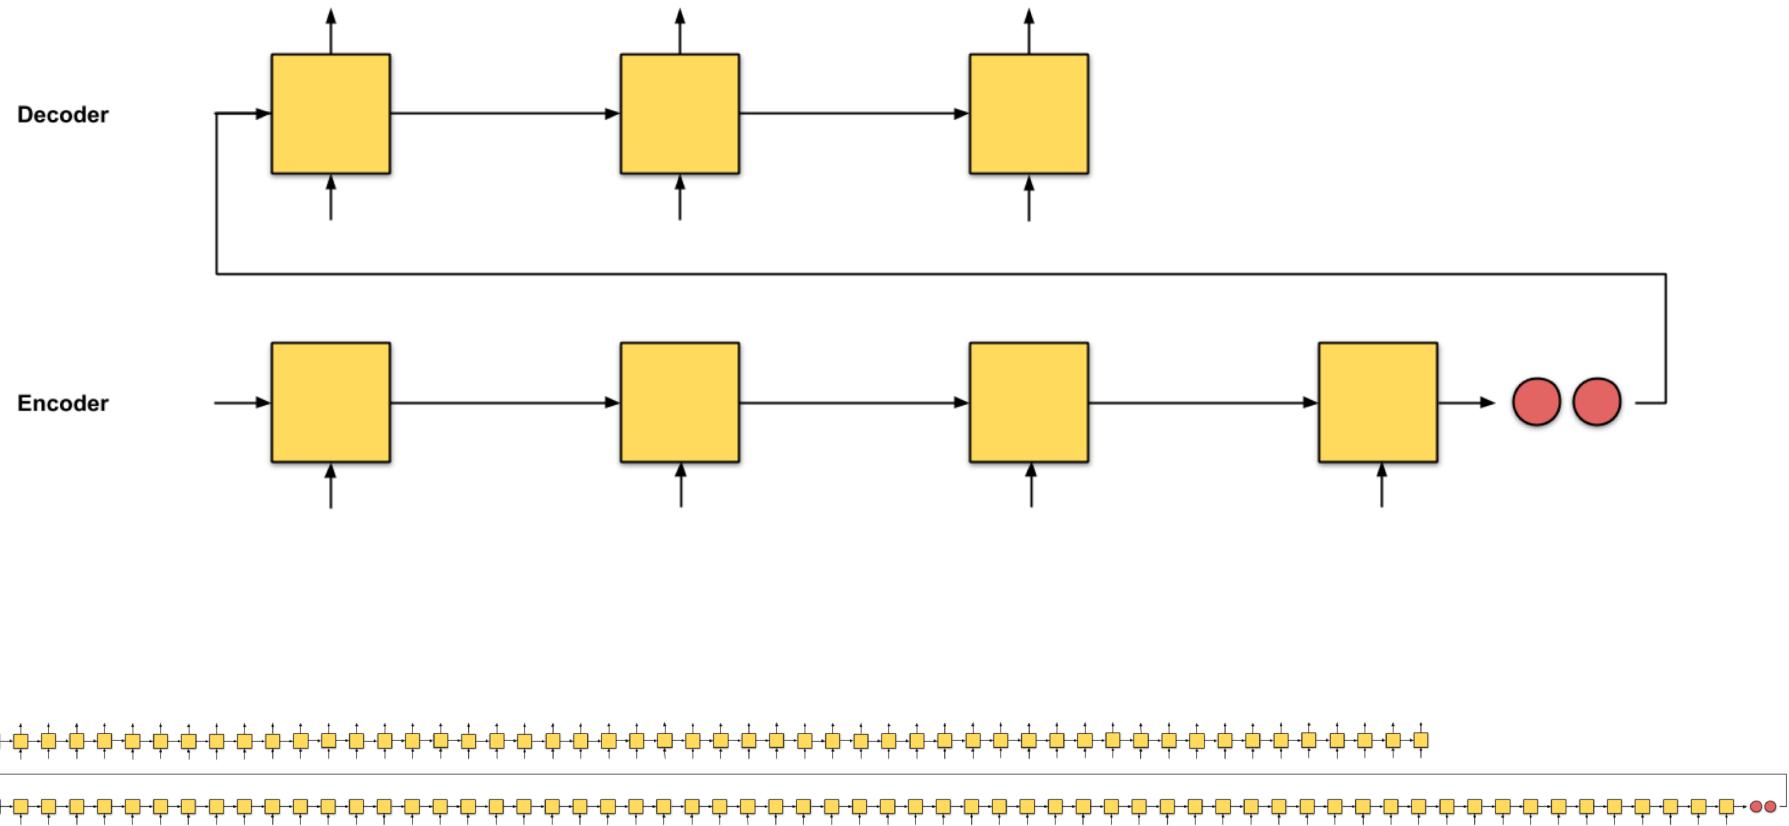


Bidirectional RNN



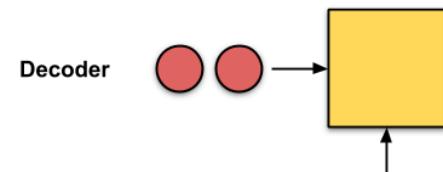
Loop forward over the sequence.
Loop backward over the sequence.
Combine output(s).

seq2seq (Sutskever et al., 2014)

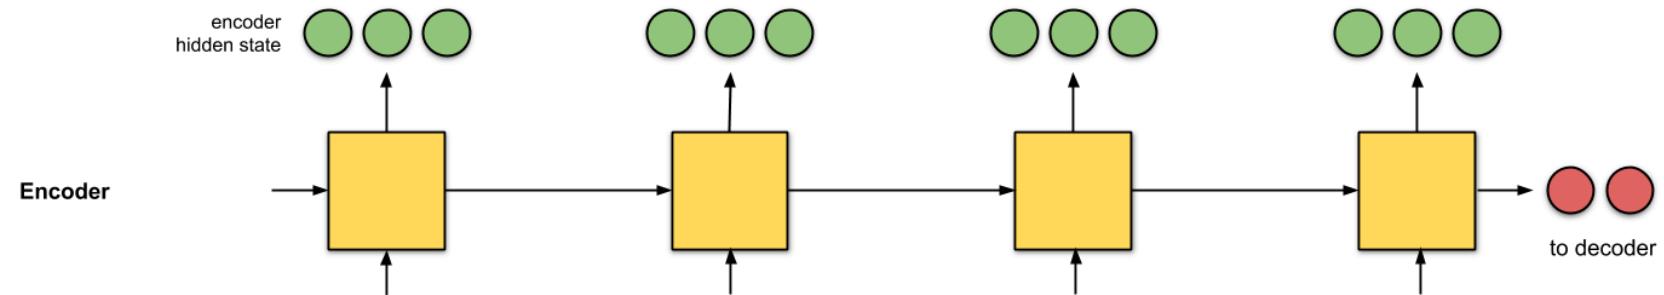


[Diagram source](#)

Attention Mechanism



Attention layer

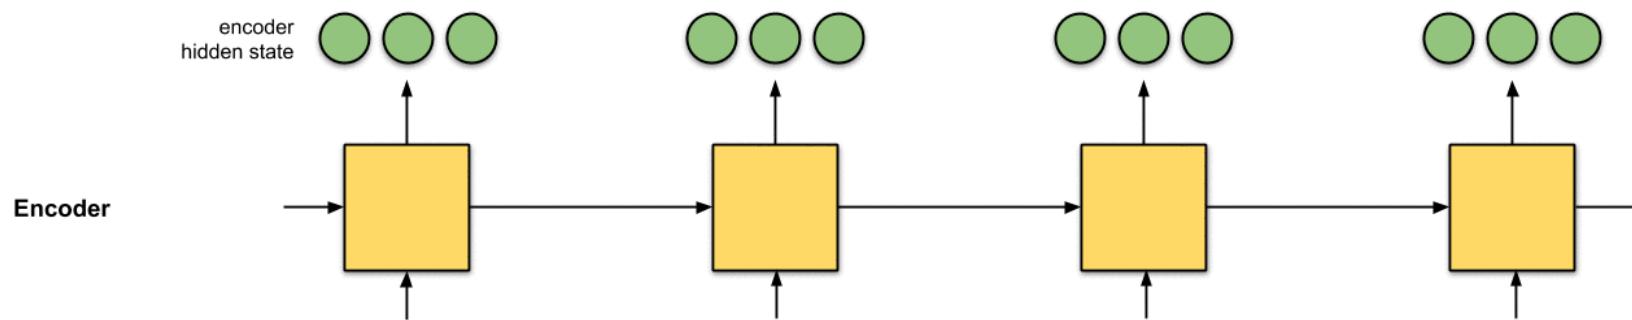
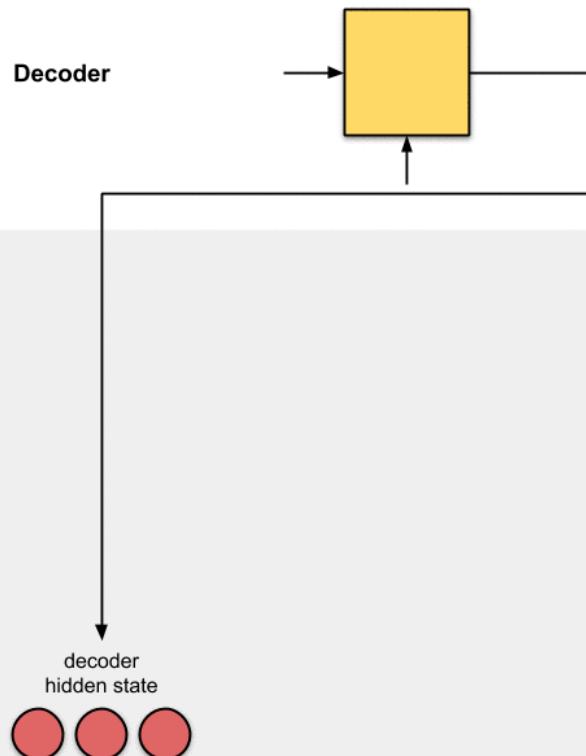


Attention Mechanism

Encoder

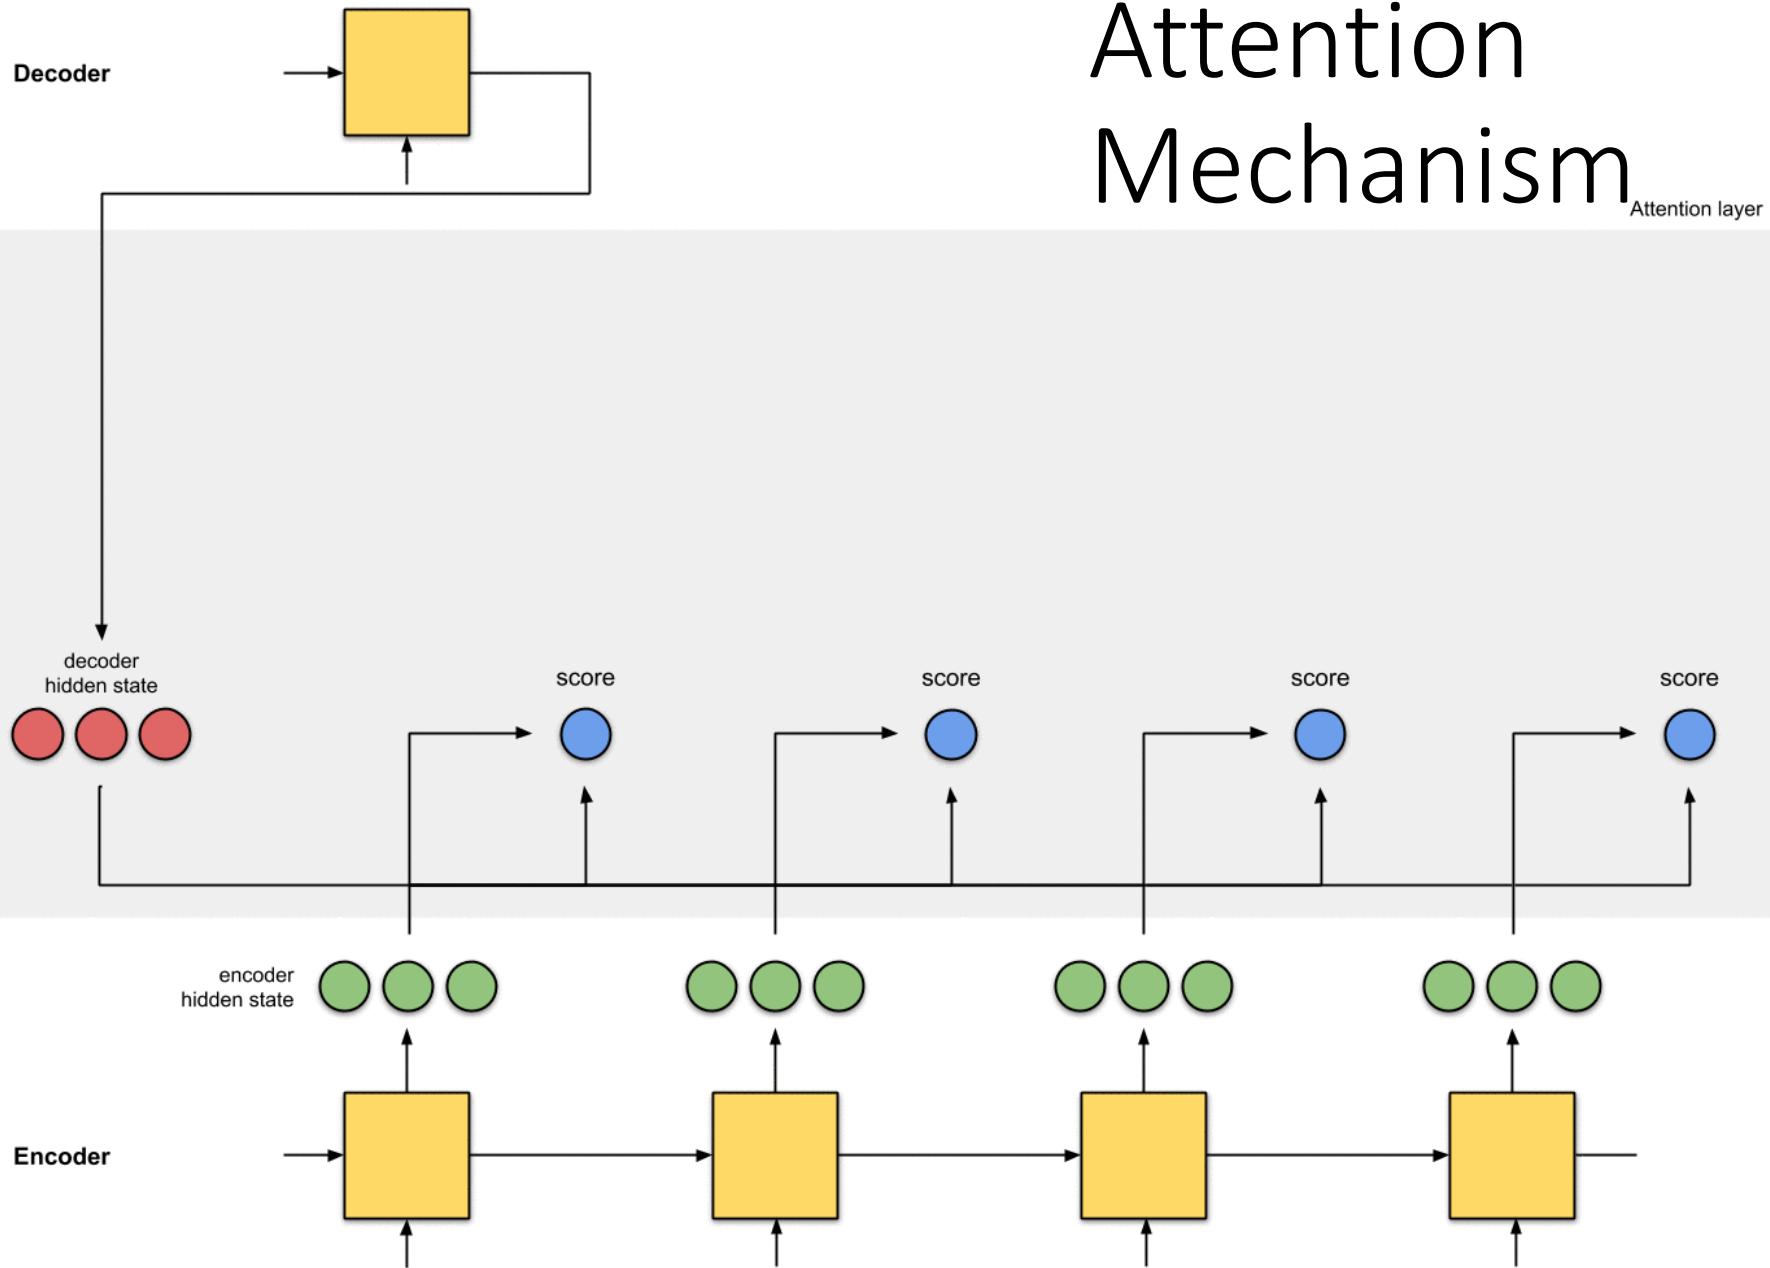
Attention Mechanism

Attention layer

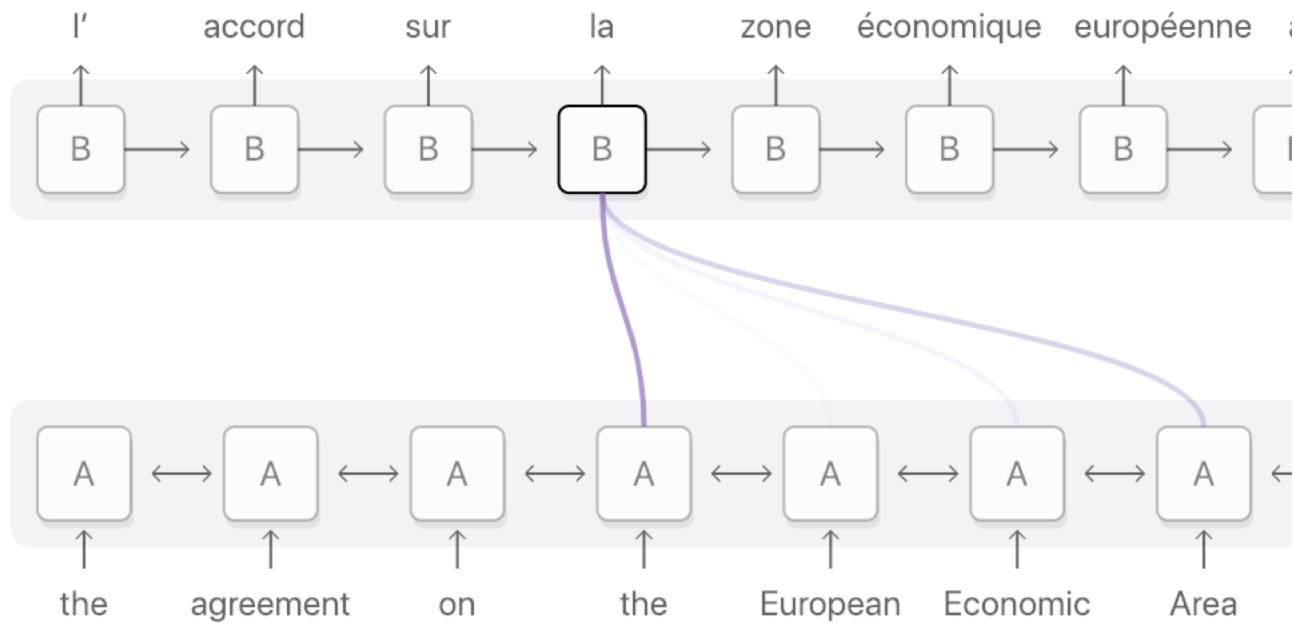


Attention Mechanism

Attention layer

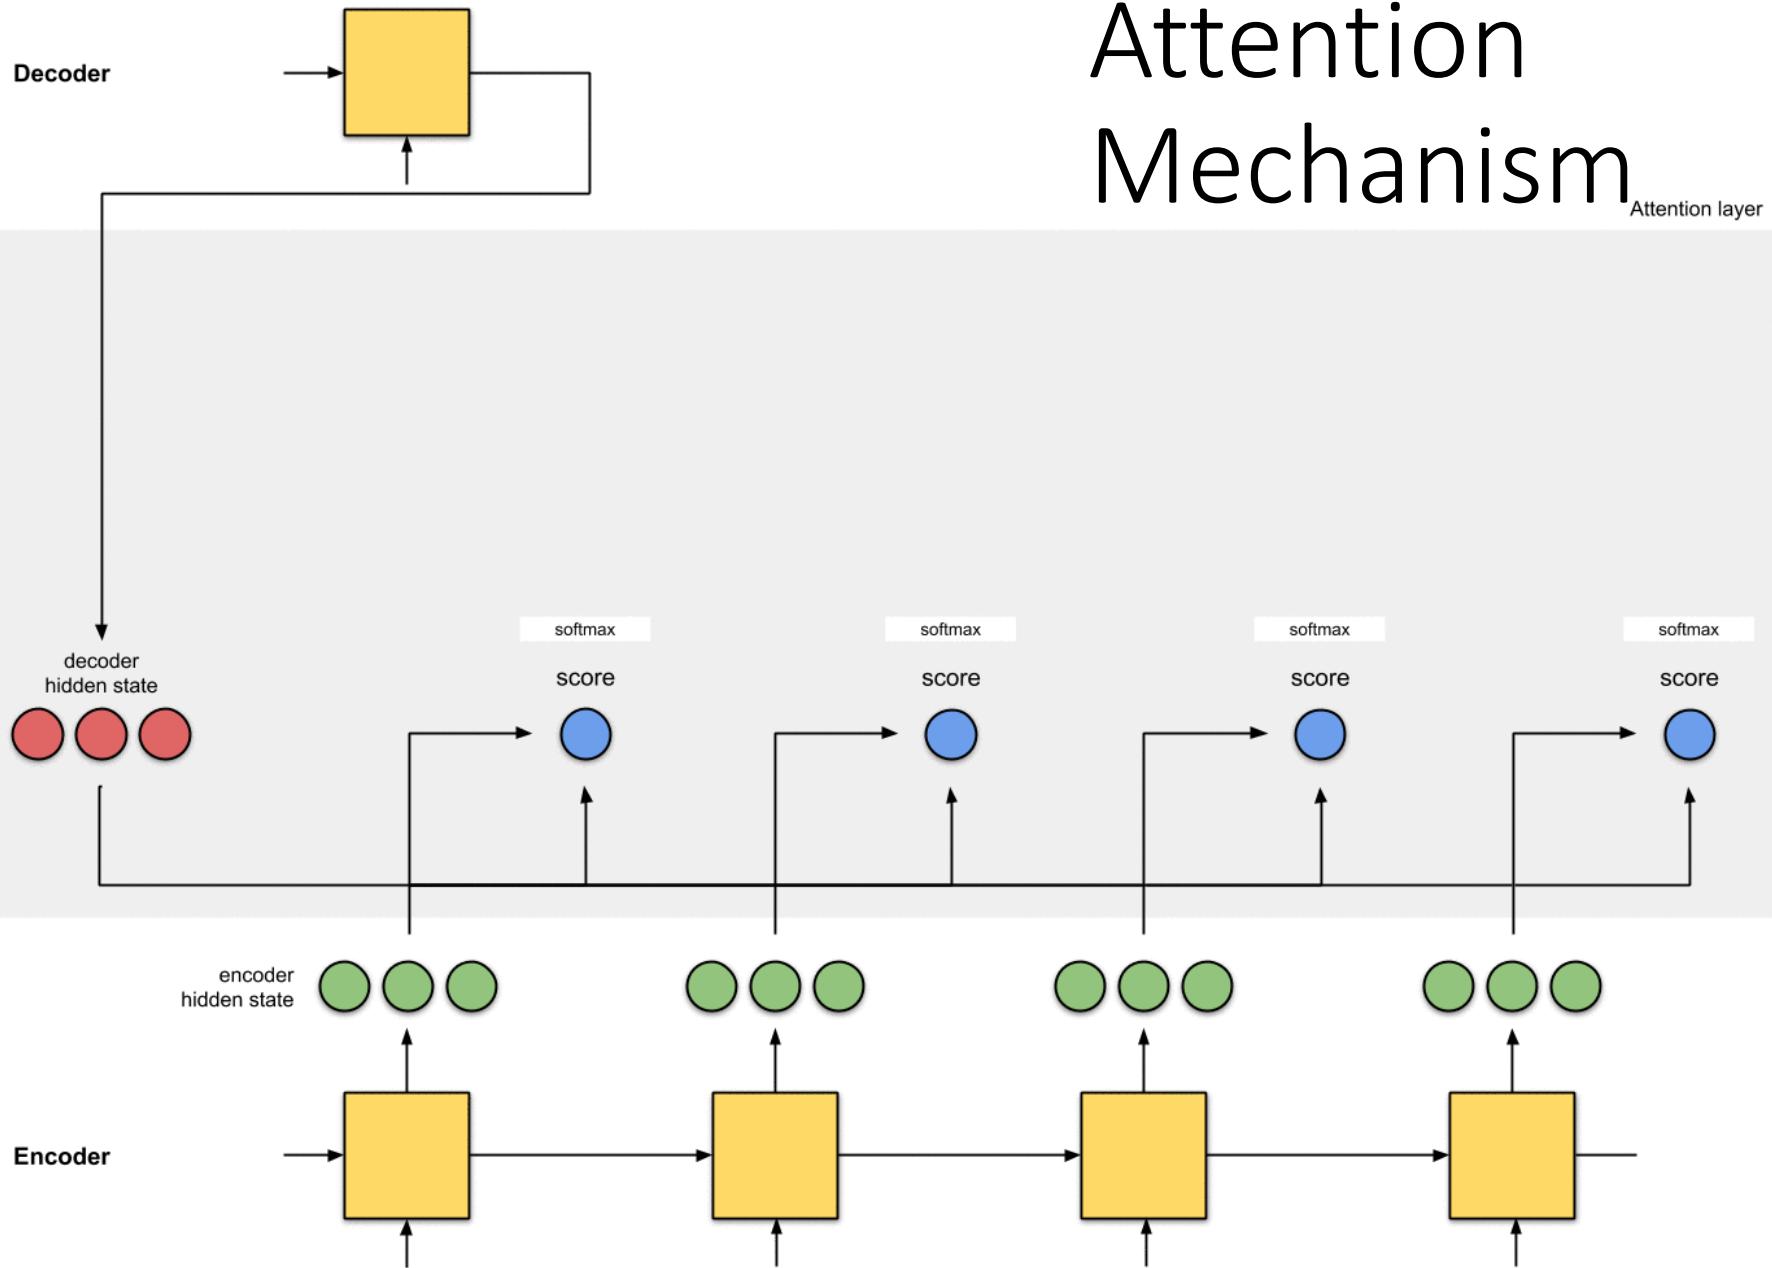


Attention Mechanism



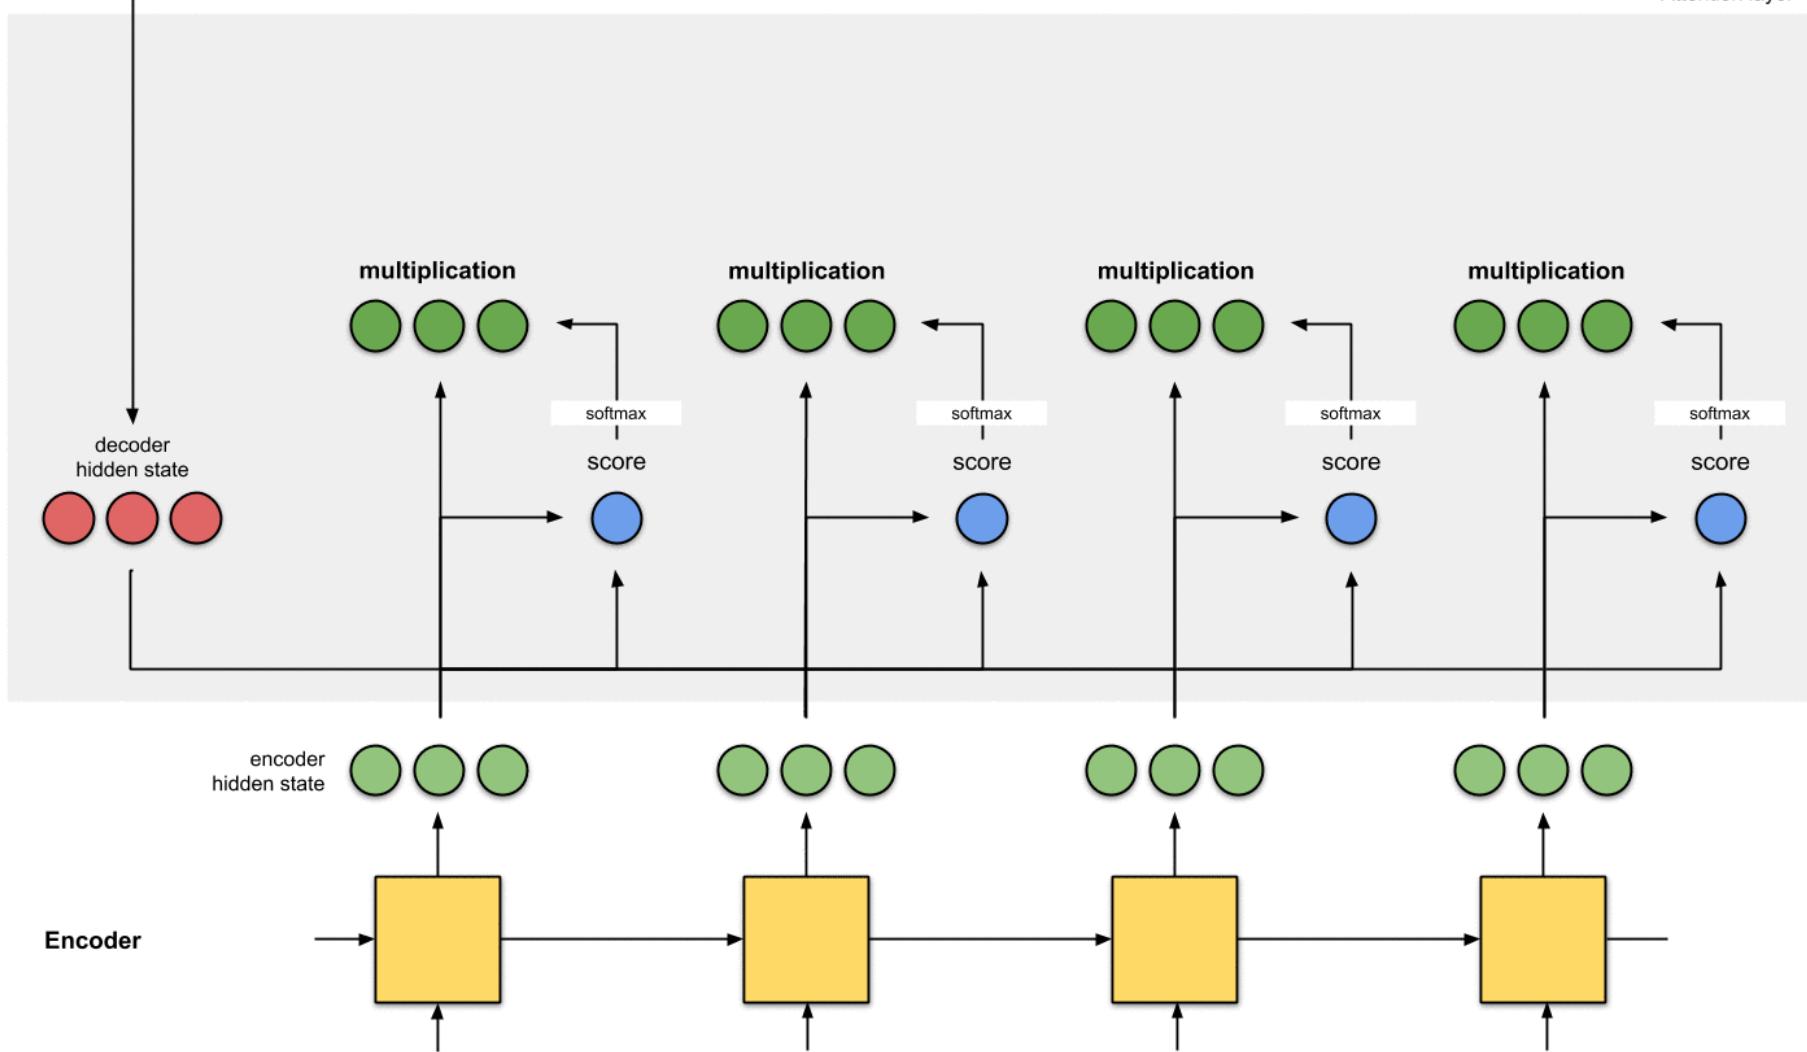
Attention Mechanism

Attention layer

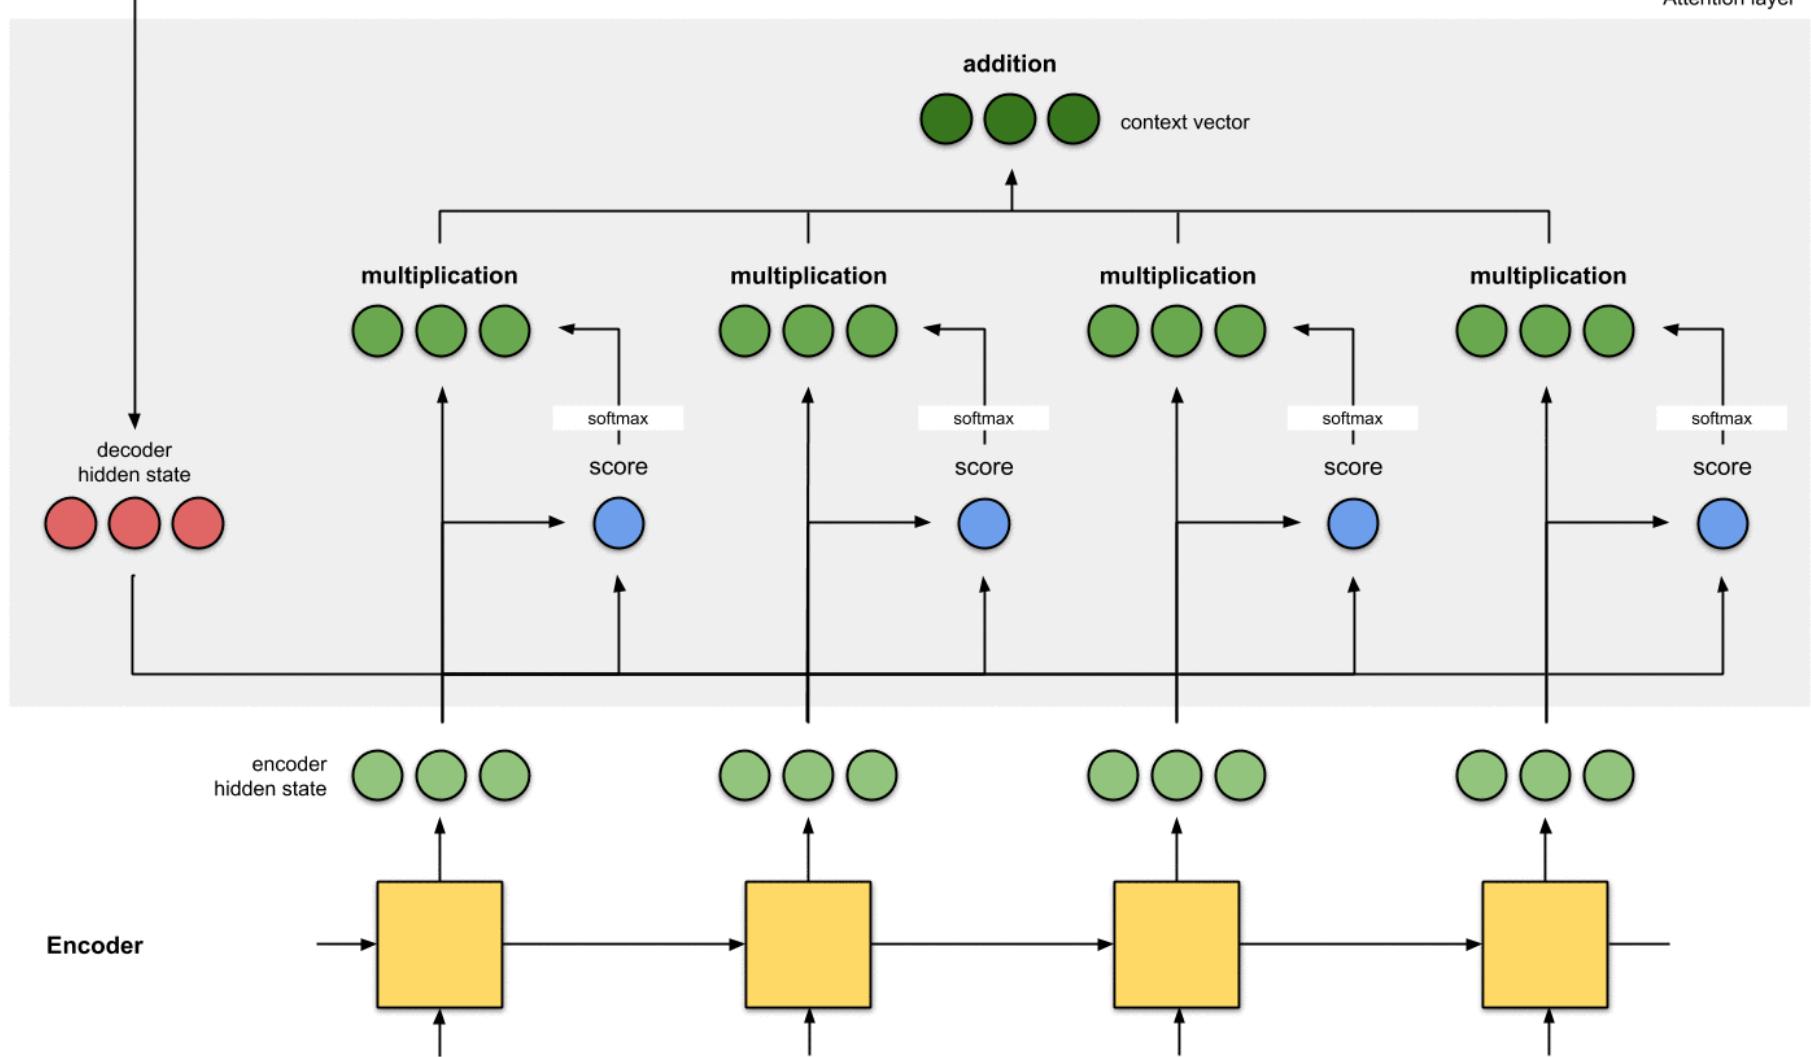


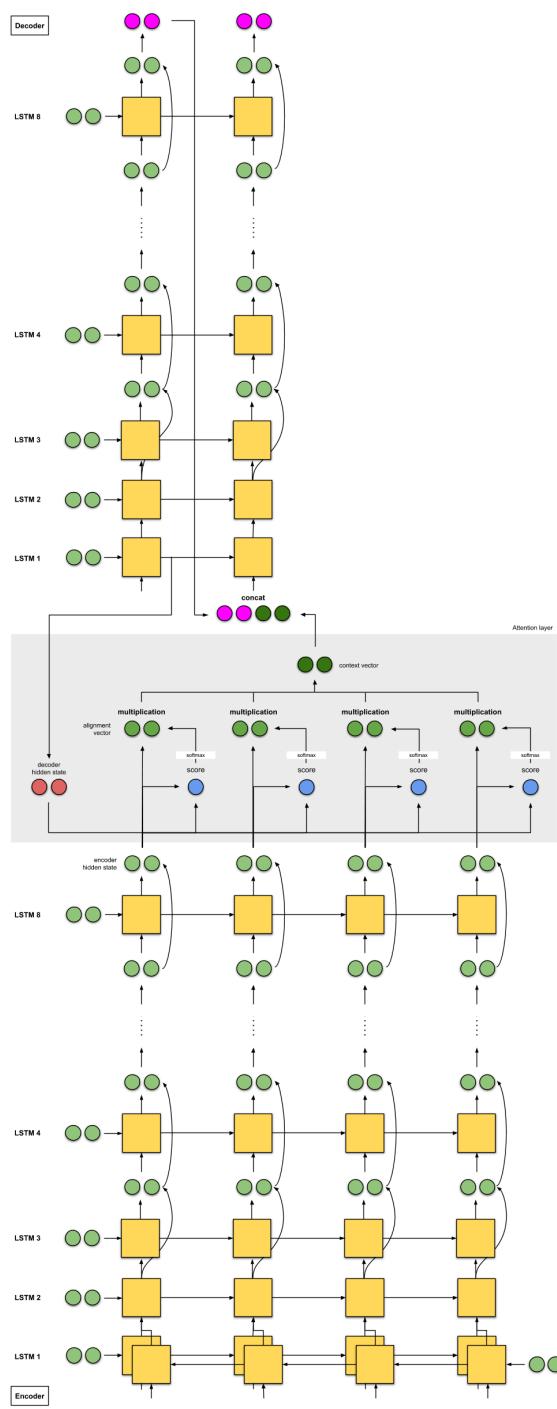
Attention Mechanism

Attention layer



Attention Mechanism





Google's Neural Machine Translation

ENCODER #2

ENCODER #1



Feed Forward
Neural Network

Feed Forward
Neural Network

z_1

z_2

Self-Attention

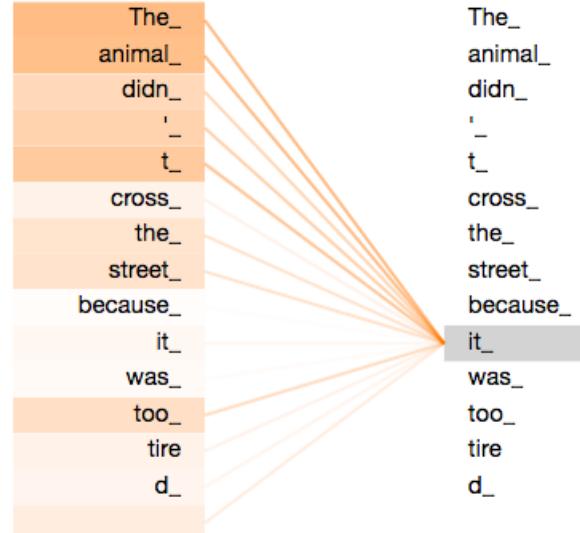
x_1

x_2

Thinking

Machines

Layer: 5 ⚡ Attention: Input - Input ⚡



From [The Illustrated Transformer](#)